

Importing Necessary Libraries

In [1]:

```
import pandas as pd
import numpy as np
import sklearn
from sklearn.model_selection import train_test_split
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import classification_report, accuracy_score
from sklearn.metrics import confusion_matrix
import numpy as np
from collections import defaultdict
import pydot
from io import StringIO
from sklearn.tree import export_graphviz
from sklearn.model_selection import GridSearchCV
from sklearn.neural_network import MLPClassifier
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.feature_selection import SelectFromModel
from sklearn.metrics import roc_auc_score
from sklearn.ensemble import VotingClassifier
from sklearn.feature_selection import RFECV
from sklearn.metrics import roc_curve
from itertools import compress
from imblearn.under_sampling import RandomUnderSampler
from imblearn.over_sampling import RandomOverSampler
import warnings
warnings.filterwarnings('ignore')

'''
TODO:

1. Try to improve
2. Desing the replace_val for each column
3. Creat preprocess procedure for every class.
'''

%matplotlib inline

rs = 101
```

Task 1. Data Selection and Distribution.

In [2]:

```
## Read Data
df = pd.read_csv("CaseStudyData.csv")
```

1 What is the proportion of cars who can be classified as a “kick”?

In [3]:

```
## Exploring the features in this dataset
print("Number of Columns: ", len(df.columns))
print("Columns: ", list(df.columns))
```

Number of Columns: 31

Columns: ['PurchaseID', 'PurchaseTimestamp', 'PurchaseDate', 'Auction', 'VehYear', 'Make', 'Color', 'Transmission', 'WheelTypeID', 'WheelType', 'VehOdo', 'Nationality', 'Size', 'TopThreeAmericanName', 'MMRAcquisitionAuctionAveragePrice', 'MMRAcquisitionAuctionCleanPrice', 'MMRAcquisitionRetailAveragePrice', 'MMRAcquisitionRetailCleanPrice', 'MMRCurrentAuctionAveragePrice', 'MMRCurrentAuctionCleanPrice', 'MMRCurrentRetailAveragePrice', 'MMRCurrentRetailCleanPrice', 'MMRCurrentRetailRatio', 'PRIMEUNIT', 'AUCGUART', 'VNST', 'VehBCost', 'IsOnlineSale', 'WarrantyCost', 'ForSale', 'IsBadBuy']

In [4]:

```
print("Number of Observations: ", len(df))
```

Number of Observations: 41476

In [5]:

```
proportionOfKicks = len(df[df['IsBadBuy'] == 1]) / len(list(df['IsBadBuy']))
print("The proportion of kicks: ", proportionOfKicks)
```

The proportion of kicks: 0.1294965763333012

2. Did you have to fix any data quality problems? Detail them.

In [6]:

```
#### PREPROCESSING STRATEGY
NEW_STRATEGY = True
ResamplingMethod = 'ros' #['ros', 'rus']
if NEW_STRATEGY:
    print("Using New Preprocessing Strategy")
    using_cat = True
    categorial_cols = ['Auction', 'VehYear', 'Make', 'Color', 'Transmission', 'WheelTypeID', 'WheelType', 'Nationality', 'Size', 'TopThreeAmericanName', 'PRIMEUNIT', 'AUCGUART', 'VNST', 'IsOnlineSale', 'ForSale' ] # Replaced by the most common
    interval_cols = ['VehOdo', 'MMRAcquisitionAuctionAveragePrice', 'MMRAcquisitionAuctionCleanPrice', 'MMRAcquisitionRetailAveragePrice', 'MMRAcquisitionRetailCleanPrice', 'VehBCost', 'WarrantyCost' ]
    drop_cols = ['PurchaseID', 'PurchaseDate', 'PurchaseTimestamp']
    questionMark_data = ['MMRCurrentAuctionAveragePrice', 'MMRCurrentAuctionCleanPrice', 'MMRCurrentRetailAveragePrice', 'MMRCurrentRetailCleanPrice', 'MMRCurrentRetailRatio']
    replaced_vals = ['?', '#VALUE!']
    if using_cat:
        categorial_cols += questionMark_data
        print("See [MMRCurrentAuctionAveragePrice" +
              "MMRCurrentAuctionCleanPrice, MMRCurrentRetailAveragePrice," +
              " MMRCurrentRetailCleanPrice, MMRCurrentRetailRatio] as Categorical
Data")
    else:
        interval_cols += questionMark_data
        print("See [MMRCurrentAuctionAveragePrice" +
              "MMRCurrentAuctionCleanPrice, MMRCurrentRetailAveragePrice," +
              " MMRCurrentRetailCleanPrice, MMRCurrentRetailRatio] as Interval Data")
    else:
        print("Using Old Preprocessing Strategy")
        drop_cols = ['PurchaseID', 'PurchaseDate']
        categorial_cols = ['Auction', 'VehYear', 'Make', 'Color', 'Transmission', 'WheelTypeID', 'WheelType', 'Nationality', 'Size', 'TopThreeAmericanName', 'PRIMEUNIT', 'AUCGUART', 'VNST', 'IsOnlineSale', 'ForSale' ] # Replaced by the most common
        interval_cols = ['PurchaseTimestamp', 'VehOdo', 'MMRAcquisitionAuctionAveragePrice', 'MMRAcquisitionAuctionCleanPrice', 'MMRAcquisitionRetailAveragePrice', 'MMRAcquisitionRetailCleanPrice', 'MMRCurrentAuctionAveragePrice', 'MMRCurrentAuctionCleanPrice', 'MMRCurrentRetailAveragePrice', 'MMRCurrentRetailCleanPrice', 'MMRCurrentRetailRatio', 'VehBCost', 'WarrantyCost' ] # Replaced by the mean
        replaced_vals = ['?', '#VALUE!']

print("Total null before Replacing: ", df.isnull().sum().sum())
```

Using New Preprocessing Strategy

See [MMRCurrentAuctionAveragePriceMMRCurrentAuctionCleanPrice, MMRCurrentRetailAveragePrice, MMRCurrentRetailCleanPrice, MMRCurrentRetailRatio] as Categorical Data

Total null before Replacing: 1691

In [7]:

```

def printColumnInfo():
    '''
    Display the information of this Dataframe
    '''

    for colName in df.columns:
        print("===== " + str(colName) + " =====")
        print("----- FIRST FIVE -----")
        print(df[colName][:5])
        print("----- DESCRIBE -----")
        print(df[colName].describe())
        print("----- COUNTS -----")
        commonList = list(df[colName].value_counts().keys())
        if len(commonList) > 100:
            print("Five Most Common: ", commonList[:5])
        else:
            print("Count List: \n", df[colName].value_counts())
        print("Num of NULL: ", df[colName].isnull().sum())
        for rep in replaced_vals:
            print("Number of "+str(rep)+" : " + str(len(df[df[colName] == rep
])))
printColumnInfo()

```

```
===== PurchaseID =====
----- FIRST FIVE -----
0      0
1      1
2      2
3      3
4      4
```

Name: PurchaseID, dtype: int64

```
----- DESCRIBE -----
count      41476.000000
mean       20737.500000
std        11973.234219
min         0.000000
25%        10368.750000
50%        20737.500000
75%        31106.250000
max        41475.000000
```

Name: PurchaseID, dtype: float64

```
----- COUNTS -----
Five Most Common: [2047, 11567, 15693, 13644, 3403]
Num of NULL: 0
Number of ? : 0
Number of #VALUE! : 0
```

```
===== PurchaseTimestamp =====
----- FIRST FIVE -----
0      1253232000
1      1253232000
2      1253232000
3      1253232000
4      1253232000
```

Name: PurchaseTimestamp, dtype: int64

```
----- DESCRIBE -----
count      4.147600e+04
mean       1.262260e+09
std        1.796895e+07
min        1.231114e+09
25%        1.247530e+09
50%        1.262045e+09
75%        1.277770e+09
max        1.293667e+09
```

Name: PurchaseTimestamp, dtype: float64

```
----- COUNTS -----
Five Most Common: [1235520000, 1259020800, 1234396800, 1264032000,
1287014400]
Num of NULL: 0
Number of ? : 0
Number of #VALUE! : 0
```

```
===== PurchaseDate =====
----- FIRST FIVE -----
0      18/09/2009 10:00
1      18/09/2009 10:00
2      18/09/2009 10:00
3      18/09/2009 10:00
4      18/09/2009 10:00
```

Name: PurchaseDate, dtype: object

```
----- DESCRIBE -----
count      41476
unique      497
top        12/02/2009 10:00
freq       242
```

Name: PurchaseDate, dtype: object

```

----- COUNTS -----
Five Most Common: ['12/02/2009 10:00', '24/11/2009 10:00', '25/02/2
009 10:00', '21/01/2010 10:00', '14/10/2010 10:00']
Num of NULL: 0
Number of ? : 0
Number of #VALUE! : 0
===== Auction =====
----- FIRST FIVE -----
0    OTHER
1    OTHER
2    OTHER
3    OTHER
4    OTHER
Name: Auction, dtype: object
----- DESCRIBE -----
count      41432
unique      3
top         MANHEIM
freq        22168
Name: Auction, dtype: object
----- COUNTS -----
Count List:
MANHEIM      22168
ADESA         11086
OTHER          8178
Name: Auction, dtype: int64
Num of NULL: 44
Number of ? : 0
Number of #VALUE! : 0
===== VehYear =====
----- FIRST FIVE -----
0    2008.0
1    2008.0
2    2008.0
3    2008.0
4    2008.0
Name: VehYear, dtype: float64
----- DESCRIBE -----
count      41432.000000
mean        2005.360615
std          1.730587
min          2001.000000
25%          2004.000000
50%          2005.000000
75%          2007.000000
max          2010.000000
Name: VehYear, dtype: float64
----- COUNTS -----
Count List:
2006.0      9630
2005.0      8682
2007.0      6514
2004.0      5792
2008.0      4177
2003.0      3554
2002.0      1879
2001.0       816
2009.0       387
2010.0        1
Name: VehYear, dtype: int64
Num of NULL: 44

```

Number of ? : 0

Number of #VALUE! : 0

===== Make =====

----- FIRST FIVE -----

0 DODGE

1 DODGE

2 CHRYSLER

3 CHEVROLET

4 DODGE

Name: Make, dtype: object

----- DESCRIBE -----

count 41432

unique 30

top CHEVROLET

freq 9548

Name: Make, dtype: object

----- COUNTS -----

Count List:

CHEVROLET 9548

DODGE 7385

FORD 6458

CHRYSLER 5259

PONTIAC 2355

KIA 1337

SATURN 1245

NISSAN 1186

JEEP 985

HYUNDAI 957

SUZUKI 842

TOYOTA 664

MITSUBISHI 569

MAZDA 532

MERCURY 527

BUICK 413

GMC 351

HONDA 263

OLDSMOBILE 146

ISUZU 82

SCION 77

VOLKSWAGEN 73

LINCOLN 54

INFINITI 27

MINI 19

ACURA 19

CADILLAC 17

SUBARU 17

LEXUS 13

VOLVO 12

Name: Make, dtype: int64

Num of NULL: 44

Number of ? : 0

Number of #VALUE! : 0

===== Color =====

----- FIRST FIVE -----

0 RED

1 RED

2 SILVER

3 RED

4 SILVER

Name: Color, dtype: object

----- DESCRIBE -----

```
count      41432
unique      17
top         SILVER
freq        8541
```

Name: Color, dtype: object

----- COUNTS -----

Count List:

```
SILVER      8541
WHITE       6890
BLUE        5855
BLACK       4392
GREY        4248
RED         3661
GOLD        3059
GREEN       1796
MAROON      1039
BEIGE       894
ORANGE      255
BROWN       249
PURPLE      205
YELLOW      141
OTHER       136
NOT AVAIL   65
?           6
```

Name: Color, dtype: int64

Num of NULL: 44

Number of ? : 6

Number of #VALUE! : 0

===== Transmission =====

----- FIRST FIVE -----

```
0    AUTO
1    AUTO
2    AUTO
3    AUTO
4    AUTO
```

Name: Transmission, dtype: object

----- DESCRIBE -----

```
count      41432
unique      4
top         AUTO
freq        39930
```

Name: Transmission, dtype: object

----- COUNTS -----

Count List:

```
AUTO       39930
MANUAL     1495
?           6
Manual      1
```

Name: Transmission, dtype: int64

Num of NULL: 44

Number of ? : 6

Number of #VALUE! : 0

===== WheelTypeID =====

----- FIRST FIVE -----

```
0    2
1    2
2    2
3    2
4    2
```

Name: WheelTypeID, dtype: object

----- DESCRIBE -----


```
count      41432
unique      5
top         1
freq       20426
```

Name: WheelTypeID, dtype: object

----- COUNTS -----

Count List:

```
1      20426
2      18791
?       1775
3        437
0         3
```

Name: WheelTypeID, dtype: int64

Num of NULL: 44

Number of ? : 1775

Number of #VALUE! : 0

===== WheelType =====

----- FIRST FIVE -----

```
0      Covers
1      Covers
2      Covers
3      Covers
4      Covers
```

Name: WheelType, dtype: object

----- DESCRIBE -----

```
count      41380
unique      4
top        Alloy
freq       20406
```

Name: WheelType, dtype: object

----- COUNTS -----

Count List:

```
Alloy      20406
Covers     18761
?          1777
Special    436
```

Name: WheelType, dtype: int64

Num of NULL: 96

Number of ? : 1777

Number of #VALUE! : 0

===== Veh0do =====

----- FIRST FIVE -----

```
0      51099.0
1      48542.0
2      46318.0
3      50413.0
4      50199.0
```

Name: Veh0do, dtype: float64

----- DESCRIBE -----

```
count      41432.000000
mean       71300.010427
std        14724.041171
min         577.000000
25%        61578.000000
50%        73128.500000
75%        82259.250000
max        480444.000000
```

Name: Veh0do, dtype: float64

----- COUNTS -----

Five Most Common: [84675.0, 85884.0, 67464.0, 72101.0, 79600.0]

Num of NULL: 44

Number of ? : 0

Number of #VALUE! : 0

===== Nationality =====

----- FIRST FIVE -----

0 AMERICAN

1 AMERICAN

2 AMERICAN

3 AMERICAN

4 AMERICAN

Name: Nationality, dtype: object

----- DESCRIBE -----

count 41432

unique 6

top AMERICAN

freq 34616

Name: Nationality, dtype: object

----- COUNTS -----

Count List:

AMERICAN 34616

OTHER ASIAN 4474

TOP LINE ASIAN 2110

USA 125

OTHER 104

? 3

Name: Nationality, dtype: int64

Num of NULL: 44

Number of ? : 3

Number of #VALUE! : 0

===== Size =====

----- FIRST FIVE -----

0 MEDIUM

1 MEDIUM

2 MEDIUM

3 COMPACT

4 MEDIUM

Name: Size, dtype: object

----- DESCRIBE -----

count 41432

unique 13

top MEDIUM

freq 17540

Name: Size, dtype: object

----- COUNTS -----

Count List:

MEDIUM 17540

LARGE 4968

MEDIUM SUV 4569

COMPACT 4035

VAN 3367

LARGE TRUCK 1897

SMALL SUV 1332

SPECIALTY 998

CROSSOVER 974

LARGE SUV 830

SMALL TRUCK 494

SPORTS 425

? 3

Name: Size, dtype: int64

Num of NULL: 44

Number of ? : 3

Number of #VALUE! : 0

```

===== TopThreeAmericanName =====
----- FIRST FIVE -----
0    CHRYSLER
1    CHRYSLER
2    CHRYSLER
3         GM
4    CHRYSLER
Name: TopThreeAmericanName, dtype: object
----- DESCRIBE -----
count      41432
unique       5
top         GM
freq       14075
Name: TopThreeAmericanName, dtype: object
----- COUNTS -----
Count List:
  GM      14075
CHRYSLER  13627
FORD      7039
OTHER     6688
?          3
Name: TopThreeAmericanName, dtype: int64
Num of NULL:  44
Number of ? : 3
Number of #VALUE! : 0
===== MMRAcquisitionAuctionAveragePrice =====
=====
----- FIRST FIVE -----
0    8566
1    8566
2    8835
3    7165
4    8566
Name: MMRAcquisitionAuctionAveragePrice, dtype: object
----- DESCRIBE -----
count      41416
unique     9271
top         0
freq       502
Name: MMRAcquisitionAuctionAveragePrice, dtype: object
----- COUNTS -----
Five Most Common:  ['0', '5480', '6311', '7811', '7644']
Num of NULL:  60
Number of ? : 7
Number of #VALUE! : 0
===== MMRAcquisitionAuctionCleanPrice =====
=====
----- FIRST FIVE -----
0    9325
1    9325
2    9428
3    7770
4    9325
Name: MMRAcquisitionAuctionCleanPrice, dtype: object
----- DESCRIBE -----
count      41429
unique    10010
top         0
freq       415
Name: MMRAcquisitionAuctionCleanPrice, dtype: object
----- COUNTS -----

```

Five Most Common: ['0', '6461', '7450', '1', '8258']

Num of NULL: 47

Number of ? : 7

Number of #VALUE! : 0

===== MMRAcquisitionRetailAveragePrice =====
=====

----- FIRST FIVE -----

0 9751

1 9751

2 10042

3 8238

4 9751

Name: MMRAcquisitionRetailAveragePrice, dtype: object

----- DESCRIBE -----

count 41429

unique 11070

top 0

freq 502

Name: MMRAcquisitionRetailAveragePrice, dtype: object

----- COUNTS -----

Five Most Common: ['0', '6418', '7316', '11114', '8756']

Num of NULL: 47

Number of ? : 7

Number of #VALUE! : 0

===== MMRAcquisitonRetailCleanPrice =====
=====

----- FIRST FIVE -----

0 10571

1 10571

2 10682

3 8892

4 10571

Name: MMRAcquisitonRetailCleanPrice, dtype: object

----- DESCRIBE -----

count 41327

unique 11583

top 0

freq 501

Name: MMRAcquisitonRetailCleanPrice, dtype: object

----- COUNTS -----

Five Most Common: ['0', '7478', '8546', '11562', '10103']

Num of NULL: 149

Number of ? : 7

Number of #VALUE! : 0

===== MMRCurrentAuctionAveragePrice =====
=====

----- FIRST FIVE -----

0 7781

1 8568

2 8137

3 7074

4 7857

Name: MMRCurrentAuctionAveragePrice, dtype: object

----- DESCRIBE -----

count 41429

unique 9183

top 0

freq 287

Name: MMRCurrentAuctionAveragePrice, dtype: object

----- COUNTS -----

Five Most Common: ['0', '?', '5480', '6311', '7269']

Num of NULL: 47
 Number of ? : 184
 Number of #VALUE! : 0

===== MMRCurrentAuctionCleanPrice =====
 =====

----- FIRST FIVE -----

0 8545
 1 9325
 2 8733
 3 7629
 4 8711

Name: MMRCurrentAuctionCleanPrice, dtype: object

----- DESCRIBE -----

count 41429
 unique 9890
 top 0
 freq 206

Name: MMRCurrentAuctionCleanPrice, dtype: object

----- COUNTS -----

Five Most Common: ['0', '?', '6461', '1', '7450']

Num of NULL: 47
 Number of ? : 184
 Number of #VALUE! : 0

===== MMRCurrentRetailAveragePrice =====
 =====

----- FIRST FIVE -----

0 11777
 1 9753
 2 9288
 3 8140
 4 8986

Name: MMRCurrentRetailAveragePrice, dtype: object

----- DESCRIBE -----

count 41409
 unique 10935
 top 0
 freq 287

Name: MMRCurrentRetailAveragePrice, dtype: object

----- COUNTS -----

Five Most Common: ['0', '?', '6418', '7316', '8756']

Num of NULL: 67
 Number of ? : 184
 Number of #VALUE! : 0

===== MMRCurrentRetailCleanPrice =====
 =====

----- FIRST FIVE -----

0 12505
 1 10571
 2 9932
 3 8739
 4 9908

Name: MMRCurrentRetailCleanPrice, dtype: object

----- DESCRIBE -----

count 41409
 unique 11363
 top 0
 freq 287

Name: MMRCurrentRetailCleanPrice, dtype: object

----- COUNTS -----

Five Most Common: ['0', '?', '7478', '8546', '10103']

Num of NULL: 67

Number of ? : 184

Number of #VALUE! : 0

===== MMRCurrentRetailRatio =====

=

----- FIRST FIVE -----

0 0.941783287

1 0.922618485

2 0.935159082

3 0.931456688

4 0.906943884

Name: MMRCurrentRetailRatio, dtype: object

----- DESCRIBE -----

count 41116

unique 25870

top #VALUE!

freq 178

Name: MMRCurrentRetailRatio, dtype: object

----- COUNTS -----

Five Most Common: ['#VALUE!', '0.858250869', '0.856073017', '0.866673265', '0.949268378']

Num of NULL: 360

Number of ? : 0

Number of #VALUE! : 178

===== PRIMEUNIT =====

----- FIRST FIVE -----

0 ?

1 ?

2 ?

3 ?

4 ?

Name: PRIMEUNIT, dtype: object

----- DESCRIBE -----

count 41432

unique 3

top ?

freq 39634

Name: PRIMEUNIT, dtype: object

----- COUNTS -----

Count List:

? 39634

NO 1764

YES 34

Name: PRIMEUNIT, dtype: int64

Num of NULL: 44

Number of ? : 39634

Number of #VALUE! : 0

===== AUCGUART =====

----- FIRST FIVE -----

0 ?

1 ?

2 ?

3 ?

4 ?

Name: AUCGUART, dtype: object

----- DESCRIBE -----

count 41432

unique 3

top ?

freq 39634

Name: AUCGUART, dtype: object

----- COUNTS -----

Count List:

? 39634

GREEN 1754

RED 44

Name: AUCGUART, dtype: int64

Num of NULL: 44

Number of ? : 39634

Number of #VALUE! : 0

===== VNST =====

----- FIRST FIVE -----

0 NC

1 NC

2 NC

3 NC

4 NC

Name: VNST, dtype: object

----- DESCRIBE -----

count 41432

unique 31

top TX

freq 9076

Name: VNST, dtype: object

----- COUNTS -----

Count List:

TX 9076

FL 5250

CO 3623

NC 3594

AZ 3383

CA 3268

OK 2595

SC 1662

TN 1471

GA 1287

VA 1093

MO 758

PA 700

NV 553

IN 486

MS 412

LA 349

NJ 317

NM 239

KY 230

AL 179

UT 165

IL 165

WV 137

OR 136

WA 136

NH 97

NE 26

OH 25

ID 14

NY 6

Name: VNST, dtype: int64

Num of NULL: 44

Number of ? : 0

Number of #VALUE! : 0

===== VehBCost =====

----- FIRST FIVE -----

```
0    7800
1    7800
2    7800
3    6000
4    7800
```

Name: VehBCost, dtype: object

----- DESCRIBE -----

```
count    41432
unique    1869
top       7500
freq      459
```

Name: VehBCost, dtype: object

----- COUNTS -----

Five Most Common: ['7500', '6500', '7800', '7200', '7000']

Num of NULL: 44

Number of ? : 29

Number of #VALUE! : 0

===== IsOnlineSale =====

----- FIRST FIVE -----

```
0    0
1    0
2    0
3    0
4    0
```

Name: IsOnlineSale, dtype: object

----- DESCRIBE -----

```
count    41432.0
unique      8.0
top        0.0
freq    31368.0
```

Name: IsOnlineSale, dtype: float64

----- COUNTS -----

Count List:

```
0.0    31368
0       8572
1.0       753
-1.0      601
1        134
?          2
4.0         1
2.0         1
```

Name: IsOnlineSale, dtype: int64

Num of NULL: 44

Number of ? : 2

Number of #VALUE! : 0

===== WarrantyCost =====

----- FIRST FIVE -----

```
0    920.0
1    834.0
2    834.0
3    671.0
4    920.0
```

Name: WarrantyCost, dtype: float64

----- DESCRIBE -----

```
count    41432.000000
mean     1273.050758
std       599.188662
min       462.000000
25%       834.000000
50%      1155.000000
75%      1623.000000
```



```

max          7498.000000
Name: WarrantyCost, dtype: float64
----- COUNTS -----
Five Most Common:  [920.0, 1974.0, 2152.0, 1215.0, 1389.0]
Num of NULL:  44
Number of ? : 0
Number of #VALUE! : 0
===== ForSale =====
----- FIRST FIVE -----
0    Yes
1    Yes
2    Yes
3    Yes
4    Yes
Name: ForSale, dtype: object
----- DESCRIBE -----
count      41476
unique       6
top        Yes
freq      27402
Name: ForSale, dtype: object
----- COUNTS -----
Count List:
  Yes      27402
YES       8544
yes       5524
?          3
No         2
0          1
Name: ForSale, dtype: int64
Num of NULL: 0
Number of ? : 3
Number of #VALUE! : 0
===== IsBadBuy =====
----- FIRST FIVE -----
0    0
1    0
2    0
3    0
4    0
Name: IsBadBuy, dtype: int64
----- DESCRIBE -----
count      41476.000000
mean        0.129497
std         0.335753
min         0.000000
25%         0.000000
50%         0.000000
75%         0.000000
max         1.000000
Name: IsBadBuy, dtype: float64
----- COUNTS -----
Count List:
  0    36105
  1    5371
Name: IsBadBuy, dtype: int64
Num of NULL: 0
Number of ? : 0
Number of #VALUE! : 0

```

In [8]:

```

if NEW_STRATEGY:

    class filling_method():
        MOST_COMMON = "MOST_COMMON"
        MEAN = "MEAN"
        CERTAIN_VALUE = "CERTAIN_VALUE"

    def replaceFunc(colName):
        for replaced, target in preprocessStrategy[colName]['replace_pairs']:
            df[colName].replace(replaced, target, inplace=True)

    def removeOutlier(colName):  # FOR THE INTERVAL ONLY
        global df
        df = df[df[colName] < df[colName].quantile(0.999)]

    def replacingValueCol(colName):
        for replaced in preprocessStrategy[colName]['replaced_vals']:
            print("In the Column: " + str(colName) + " : " + str(len(
                df[df[colName] == replaced])) + ", " + str(replaced) + "have been
replaced by null")
            # Replacing the null in this process #Inplacing for saving the memory
            df[colName].replace(replaced, float('nan'), inplace=True)

    def loweringCol(colName):
        df[colName] = df[colName].str.lower()

    def fillingTheNullValue(colName):  # method can be ["MEAN", "MOST_COMMON"]
        if preprocessStrategy[colName]['filling_method'] == filling_method.MEAN:
            df[colName] = df[colName].astype('float')
            df[colName].fillna(df[colName].astype(
                'float').mean(), inplace=True)
        elif preprocessStrategy[colName]['filling_method'] == filling_method.MOST_COMMON:
            df[colName] = df[colName].astype('category')
            df[colName].fillna(df[colName].astype(
                'category').describe()['top'], inplace=True)
        elif preprocessStrategy[colName]['filling_method'] == filling_method.CERTAIN_VALUE:
            df[colName] = df[colName].astype('category')
            df[colName] = df[colName].cat.add_categories(
                preprocessStrategy[colName]['filling_value'])
            df[colName].fillna(preprocessStrategy[colName]
                               ['filling_value'], inplace=True)

    def filterOutRareValue(colName):

        def checkingKeepValue(v, savingValues):
            if v in savingValues:
                return v
            return "LESS_FREQ"

        k = [v for v in df[colName].value_counts().values if v >
              preprocessStrategy[colName]['min_freq']]
        savingValues = df[colName].value_counts().keys()[:len(k)]

        df[colName] = [checkingKeepValue(v, savingValues) for v in df[colName]]

```

```

def changeToType(colName):
    df[colName] = df[colName].astype(
        preprocessStrategy[colName]['changeToType'])

def newData_prep(df):
    """
    For Preprocessing through the whole dictionary
    """
    df.drop(drop_cols, axis=1, inplace=True)

    for colName in df.columns: # df.columns:

        print("Preprocess the col: " + colName)

        for stra in preprocessStrategy[colName]['strategies']:
            if not stra:
                continue
            stra(colName)

    if not using_cat:
        df['MMRCurrentRetailRatio'] = df['MMRCurrentRetailAveragePrice'] / \
            (df['MMRCurrentRetailCleanPrice']+1e-8) # Prvent divided by 0

    return df

preprocessStrategy = defaultdict(dict)

preprocessStrategy['Auction'] = {
    "strategies":
        [
            replacingValueCol,
            loweringCol,
            fillingTheNullValue,
        ],
    "replaced_vals": ['?'],
    "filling_method": filling_method.MOST_COMMON
}

preprocessStrategy['VehYear'] = {
    "strategies":
        [
            fillingTheNullValue,
        ],
    "filling_method": filling_method.CERTAIN_VALUE,
    "filling_value": "UNKNOWN_VALUE"
}

preprocessStrategy['Make'] = {
    "strategies":
        [
            loweringCol,
            fillingTheNullValue,
        ],
    "filling_method": filling_method.MOST_COMMON
}

preprocessStrategy['Color'] = {
    "strategies":
        [
            loweringCol,
            replacingValueCol,

```

```

        fillingTheNullValue,
    ],
    "replaced_vals": ['?'],
    "filling_method": filling_method.MOST_COMMON
}

preprocessStrategy['Transmission'] = {
    "strategies":
    [
        loweringCol,
        replacingValueCol,
        fillingTheNullValue,
    ],
    "replaced_vals": ['?'],
    "filling_method": filling_method.MOST_COMMON
}

preprocessStrategy['WheelTypeID'] = {
    "strategies":
    [
        fillingTheNullValue,
    ],
    "filling_method": filling_method.MOST_COMMON
}

preprocessStrategy['WheelType'] = {
    "strategies":
    [
        loweringCol,
        fillingTheNullValue,
    ],
    "filling_method": filling_method.MOST_COMMON
}

preprocessStrategy['Veh0do'] = {
    "strategies":
    [
        fillingTheNullValue,
    ],
    "filling_method": filling_method.MEAN
}

preprocessStrategy['Nationality'] = { # Should I merge USA with AMERICAN?
    "strategies":
    [
        replaceFunc,
        loweringCol,
        replacingValueCol,
        fillingTheNullValue,
    ],
    "replaced_vals": ['?'],
    "filling_method": filling_method.MOST_COMMON,
    "replace_pairs": [("USA", "AMERICAN")]
}

preprocessStrategy['Size'] = {
    "strategies":
    [
        loweringCol,
        replacingValueCol,

```

```

        fillingTheNullValue,
    ],
    "replaced_vals": ['?'],
    "filling_method": filling_method.MOST_COMMON
}

preprocessStrategy['TopThreeAmericanName'] = {
    "strategies":
        [
            loweringCol,
            replacingValueCol,
            fillingTheNullValue,
        ],
    "replaced_vals": ['?'],
    "filling_method": filling_method.MOST_COMMON
}

preprocessStrategy['TopThreeAmericanName'] = {
    "strategies":
        [
            loweringCol,
            replacingValueCol,
            fillingTheNullValue,
        ],
    "replaced_vals": ['?'],
    "filling_method": filling_method.MOST_COMMON
}

preprocessStrategy['MMRAcquisitionAuctionAveragePrice'] = {
    "strategies":
        [
            replacingValueCol,
            fillingTheNullValue,
        ],
    "replaced_vals": ['?'],
    "filling_method": filling_method.MEAN
}

preprocessStrategy['MMRAcquisitionAuctionCleanPrice'] = {
    "strategies":
        [
            replacingValueCol,
            fillingTheNullValue,
        ],
    "replaced_vals": ['?'],
    "filling_method": filling_method.MEAN
}

preprocessStrategy['MMRAcquisitionRetailAveragePrice'] = {
    "strategies":
        [
            replacingValueCol,
            fillingTheNullValue,
        ],
    "replaced_vals": ['?'],
    "filling_method": filling_method.MEAN
}

preprocessStrategy['MMRAcquisitonRetailCleanPrice'] = {
    "strategies":

```

```

        [
            replacingValueCol,
            fillingTheNullValue,
        ],
        "replaced_vals": ['?'],
        "filling_method": filling_method.MEAN
    }

#####

int_stra = {
    "strategies":
        [
            replacingValueCol,
            fillingTheNullValue,
        ],
        "replaced_vals": ['?', '#VALUE!'], # GOT 184 '?'
        "filling_method": filling_method.MEAN,
    }

cat_stra = { # HOW DO WE DEAL WITH ? in this column
    "strategies":
        [
            filterOutRareValue,
            fillingTheNullValue,
        ],
        # "replaced_vals": ['?'], # GOT 184 '?'
        "filling_method": filling_method.CERTAIN_VALUE,
        "filling_value": 'NULL',
        "min_freq": 50
    }

preprocessStrategy['MMRCurrentAuctionAveragePrice'] \
= preprocessStrategy['MMRCurrentAuctionCleanPrice'] \
= preprocessStrategy['MMRCurrentRetailAveragePrice'] \
= preprocessStrategy['MMRCurrentRetailCleanPrice'] \
= preprocessStrategy['MMRCurrentRetailRatio'] \
= cat_stra if using_cat else int_stra

#####

preprocessStrategy['PRIMEUNIT'] = { # HOW DO WE DEAL WITH ? in this column
    "strategies":
        [
            loweringCol,
            fillingTheNullValue,
        ],
        # "replaced_vals": ['?'], # GOT 184 '?'
        "filling_method": filling_method.CERTAIN_VALUE,
        "filling_value": 'NULL',
    }

preprocessStrategy['AUCGUART'] = { # HOW DO WE DEAL WITH ? in this column
    "strategies":
        [
            loweringCol,
            fillingTheNullValue,
        ],
        # "replaced_vals": ['?'], # GOT 184 '?'
        "filling_method": filling_method.CERTAIN_VALUE,
        "filling_value": 'NULL',
    }

```

```

    }

    preprocessStrategy['VNST'] = { # HOW DO WE DEAL WITH ? in this column
        "strategies":
            [
                loweringCol,
                fillingTheNullValue,
            ],
        # "replaced_vals": ['?'], # GOT 184 '?'
        "filling_method": filling_method.CERTAIN_VALUE,
        "filling_value": 'NULL',
    }

    preprocessStrategy['VehBCost'] = { # HOW DO WE DEAL WITH ? in this column
        "strategies":
            [
                replacingValueCol,
                fillingTheNullValue,
            ],
        "replaced_vals": ['?'], # GOT 184 '?'
        "filling_method": filling_method.MEAN
    }

    preprocessStrategy['IsOnlineSale'] = { # HOW DO WE DEAL WITH ? in this column
mn
        "strategies":
            [
                replacingValueCol,
                changeToType,
                fillingTheNullValue,
            ],
        "replaced_vals": ['?', 2.0, 4.0], # GOT 184 '?'
        "filling_method": filling_method.MOST_COMMON,
        "changeToType": 'float'
    }

    preprocessStrategy['WarrantyCost'] = { # HOW DO WE DEAL WITH ? in this column
mn
        "strategies":
            [
                fillingTheNullValue,
            ],
        "replaced_vals": ['?'], # GOT 184 '?'
        "filling_method": filling_method.MEAN,
    }

    preprocessStrategy['ForSale'] = { # HOW DO WE DEAL WITH ? in this column
        "strategies":
            [
                loweringCol,
                replacingValueCol,
                fillingTheNullValue,
            ],
        "replaced_vals": ['?', 0], # GOT 184 '?'
        "filling_method": filling_method.MOST_COMMON,
    }

    # HOW DO WE DEAL WITH ? in this column
    preprocessStrategy['IsBadBuy'] = {"strategies": [None]}

```

```

newData_prep(df)

else:

    def data_prep(df):
        '''
        For Preprocessing the Data (OLD_METHOD)
        '''

        # Check the replaced values are not in the dataset

        for colName in df.columns:

            if colName in categorial_cols:

                if colName == "IsOnlineSale":
                    df[colName] = df[colName].astype(
                        'float').astype('category')
                    df[colName].fillna(df[colName].astype(
                        'category').describe()['top'], inplace=True)

                # Try to lower the data if the data type is string
                try:
                    df[colName] = df[colName].str.lower()
                except:
                    print(colName, " can't be lowered")

                for replaced in replaced_vals:
                    print("In the Column: " + str(colName) + ": " +
                        str(len(df[df[colName] == replaced])) + " -> " + str(r
eplaced))
                    df[colName].replace(replaced, float('nan'), inplace=True)

                df[colName] = df[colName].astype('category')

                # Replacing the null by the most common category
                df[colName].fillna(df[colName].astype(
                    'category').describe()['top'], inplace=True)

            if colName in interval_cols:

                if colName == "MMRCurrentRetailRatio": # Dealing with this calc
ulated value at the last
                    continue

                for replaced in replaced_vals:
                    print("In the Column: " + str(colName) + ": " +
                        str(len(df[df[colName] == replaced))) + " -> " + str(r
eplaced))
                    df[colName].replace(replaced, float('nan'), inplace=True)

                df[colName] = df[colName].astype('float')

                # Removing outlier
                df = df[df[colName] < df[colName].quantile(0.999)]

                # Replacing the null by the mean
                df[colName].fillna(df[colName].astype(
                    'float').mean(), inplace=True)

```



```

df['MMRCurrentRetailRatio'] = df['MMRCurrentRetailAveragePrice'] / \
    (df['MMRCurrentRetailCleanPrice']+1e-8) # Prvent divided by 0

df.drop(drop_cols, axis=1, inplace=True)

return df

df = data_prep(df)

```

Preprocess the col: Auction

In the Column: Auction : 0, ?have been replaced by null

Preprocess the col: VehYear

Preprocess the col: Make

Preprocess the col: Color

In the Column: Color : 6, ?have been replaced by null

Preprocess the col: Transmission

In the Column: Transmission : 6, ?have been replaced by null

Preprocess the col: WheelTypeID

Preprocess the col: WheelType

Preprocess the col: VehOdo

Preprocess the col: Nationality

In the Column: Nationality : 3, ?have been replaced by null

Preprocess the col: Size

In the Column: Size : 3, ?have been replaced by null

Preprocess the col: TopThreeAmericanName

In the Column: TopThreeAmericanName : 3, ?have been replaced by null

Preprocess the col: MMRAcquisitionAuctionAveragePrice

In the Column: MMRAcquisitionAuctionAveragePrice : 7, ?have been replaced by null

Preprocess the col: MMRAcquisitionAuctionCleanPrice

In the Column: MMRAcquisitionAuctionCleanPrice : 7, ?have been replaced by null

Preprocess the col: MMRAcquisitionRetailAveragePrice

In the Column: MMRAcquisitionRetailAveragePrice : 7, ?have been replaced by null

Preprocess the col: MMRAcquisitonRetailCleanPrice

In the Column: MMRAcquisitonRetailCleanPrice : 7, ?have been replaced by null

Preprocess the col: MMRCurrentAuctionAveragePrice

Preprocess the col: MMRCurrentAuctionCleanPrice

Preprocess the col: MMRCurrentRetailAveragePrice

Preprocess the col: MMRCurrentRetailCleanPrice

Preprocess the col: MMRCurrentRetailRatio

Preprocess the col: PRIMEUNIT

Preprocess the col: AUCGUART

Preprocess the col: VNST

Preprocess the col: VehBCost

In the Column: VehBCost : 29, ?have been replaced by null

Preprocess the col: IsOnlineSale

In the Column: IsOnlineSale : 2, ?have been replaced by null

In the Column: IsOnlineSale : 1, 2.0have been replaced by null

In the Column: IsOnlineSale : 1, 4.0have been replaced by null

Preprocess the col: WarrantyCost

Preprocess the col: ForSale

In the Column: ForSale : 3, ?have been replaced by null

In the Column: ForSale : 0, 0have been replaced by null

Preprocess the col: IsBadBuy




3. Can you identify any clear patterns by initial exploration of the data using histogram or box plot?




In [9]:




```
def plotAllCols (df):  
    for colName in df.columns:  
        plt.figure(figsize=(20,10))  
        if colName in categorical_cols:  
            ### if it's categorical column, plot hist diagram  
            sns.countplot(x=colName, data = df, hue="IsBadBuy")  
        elif colName in interval_cols:  
            ### if it's interval column, plot box diagram  
            sns.boxplot(x="IsBadBuy", y=colName, data = df )
```




In [10]:




```
plotAllCols(df)
```





























<Figure size 1440x720 with 0 Axes>

4. What variables did you include in the analysis and what were their roles and measurement level set? Justify your choice

In []:

5. What distribution scheme did you use? What data partitioning allocation did you set? Explain your selection.

In [11]:

```
# Change to the dummy
df = pd.get_dummies(df)

feature_names = df.drop("IsBadBuy", axis=1).columns
print("Num of Features:", len(feature_names))

### Split to the training and test set.
# The test size is 3%

# y = df['IsBadBuy']
# X = df.drop(['IsBadBuy'], axis=1)
# X_mat = X.as_matrix()

# X_train, X_test, y_train, y_test = train_test_split(X_mat, y, test_size=0.3, s
tratify=y, random_state=rs)

X_train, X_test, y_train, y_test = train_test_split(df.drop("IsBadBuy", axis=1),
df['IsBadBuy'], test_size=0.3, stratify=df['IsBadBuy'], random_state=rs)

if ResamplingMethod == 'ros':
    print("Using ROS Resmapling")
    ros = RandomOverSampler(random_state=rs)
    X_train, y_train = ros.fit_resample(X_train, y_train)
elif ResamplingMethod == 'rus':
    print("Using RUS Resmapling")
    rus = RandomUnderSampler(random_state=rs)
    X_train, y_train = rus.fit_resample(X_train, y_train)
else:
    print("No Resampling Method Used")
```

Num of Features: 198
Using ROS Resmapling

In [12]:

```
print("Number of Training: ", len(X_train))
print("Number of Test: ", len(X_test) )
```

Number of Training: 50546
Number of Test: 12443

Task 2. Predictive Modeling Using Decision Trees

1. Python: Build a decision tree using the default setting.

In [13]:

```
def printLRTopImportant(model, top = 5):

    coef = model.coef_[0]
    indices = np.argsort(np.absolute(coef))
    indices = np.flip(indices, axis=0)
    indices = indices[:top]
    for i in indices:
        print(feature_names[i], ': ', coef[i])

def analyse_feature_importance(dm_model, feature_names, n_to_display=20):
    # grab feature importances from the model
    importances = dm_model.feature_importances_

    # sort them out in descending order
    indices = np.argsort(importances)
    indices = np.flip(indices, axis=0)

    # limit to 20 features, you can leave this out to print out everything
    indices = indices[:n_to_display]

    for i in indices:
        print(feature_names[i], ': ', importances[i])

def visualize_decision_tree(dm_model, feature_names, save_name):
    dotfile = StringIO()
    export_graphviz(dm_model, out_file=dotfile, feature_names=feature_names)
    graph = pydot.graph_from_dot_data(dotfile.getvalue())
    graph[0].write_png(save_name) # saved in the following file
```

In [14]:

```
# simple decision tree training
model = DecisionTreeClassifier(random_state=rs)
model.fit(X_train, y_train)
```

Out[14]:

```
DecisionTreeClassifier(class_weight=None, criterion='gini', max_depth=
None,
                        max_features=None, max_leaf_nodes=None,
                        min_impurity_decrease=0.0, min_impurity_split=None,
                        min_samples_leaf=1, min_samples_split=2,
                        min_weight_fraction_leaf=0.0, presort=False, random_state=101,
                        splitter='best')
```

a. What is the classification accuracy on training and test datasets?

In [15]:

```
print("Train accuracy:", model.score(X_train, y_train))
print("Test accuracy:", model.score(X_test, y_test))
y_pred = model.predict(X_test)
print(classification_report(y_test, y_pred))
confusion_matrix(y_test, y_pred) ## Confusion Matrix on the TestSet
```

Train accuracy: 0.9994856170616864

Test accuracy: 0.8256047576950896

	precision	recall	f1-score	support
0	0.90	0.90	0.90	10832
1	0.33	0.33	0.33	1611
micro avg	0.83	0.83	0.83	12443
macro avg	0.61	0.61	0.61	12443
weighted avg	0.83	0.83	0.83	12443

Out[15]:

```
array([[9742, 1090],
       [1080, 531]])
```

b. What is the size of tree (i.e. number of nodes)?

In [16]:

```
print("Number of nodes: ", model.tree_.node_count)
```

Number of nodes: 7451

c. How many leaves are in the tree that is selected based on the validation dataset?

In []:

d. Which variable is used for the first split? What are the competing splits for this first split?

In [17]:

```
visualize_decision_tree(model, df.drop("IsBadBuy", axis=1).columns, "Tree_Struct.png")
```

e. What are the 5 important variables in building the tree?

In [18]:

```
analyse_feature_importance(model, df.drop("IsBadBuy", axis=1).columns, 5)
```

WheelTypeID_? : 0.13551426074337208

VehBCost : 0.12697398520218173

VehOdo : 0.09322697324736888

MMRAcquisitionAuctionCleanPrice : 0.0728830956066675

MMRAcquisitionRetailAveragePrice : 0.06240815240105831

f. Report if you see any evidence of model overfitting.

In []:

g. Did changing the default setting (i.e., only focus on changing the setting of the number of splits to create a node) help improving the model? Answer the above questions on the best performing tree.

In []:

2. Python: Build another decision tree tuned with GridSearchCV

In []:

In [19]:

```
# grid search CV
params = {'criterion': ['gini', 'entropy'],
          'max_depth': list(range(2,7)) + [200, 500] + list(range(1, 6000, 1000))
          + [None],
          'splitter': ['best', 'random'],
          'min_samples_leaf': range(1, 4),
          'min_samples_split': [2, 0.5, 0.3],
          'max_features': ['auto', 'sqrt', 'log2', None],
          'class_weight': ['balanced', None]
          }

cv = GridSearchCV(param_grid=params, estimator=DecisionTreeClassifier(random_state=rs), cv=3)
cv.fit(X_train, y_train)
```

Out[19]:

```
GridSearchCV(cv=3, error_score='raise-deprecating',
             estimator=DecisionTreeClassifier(class_weight=None, criterion
='gini', max_depth=None,
             max_features=None, max_leaf_nodes=None,
             min_impurity_decrease=0.0, min_impurity_split=None,
             min_samples_leaf=1, min_samples_split=2,
             min_weight_fraction_leaf=0.0, presort=False, random_state=101,
             splitter='best'),
             fit_params=None, iid='warn', n_jobs=None,
             param_grid={'criterion': ['gini', 'entropy'], 'max_depth':
[2, 3, 4, 5, 6, 200, 500, 1, 1001, 2001, 3001, 4001, 5001, None], 'splitter': ['best', 'random'], 'min_samples_leaf': range(1, 4), 'min_samples_split': [2, 0.5, 0.3], 'max_features': ['auto', 'sqrt', 'log2', None], 'class_weight': ['balanced', None]},
             pre_dispatch='2*n_jobs', refit=True, return_train_score='warn',
             scoring=None, verbose=0)
```

a. What is the classification accuracy on training and test datasets?

In [20]:

```
print("Train accuracy:", cv.score(X_train, y_train))
print("Test accuracy:", cv.score(X_test, y_test))

# test the best model
y_pred = cv.predict(X_test)
print(classification_report(y_test, y_pred))
# print parameters of the best model
print(cv.best_params_)

dt_model = cv.best_estimator_
```

Train accuracy: 0.9994856170616864

Test accuracy: 0.8243188941573576

	precision	recall	f1-score	support
0	0.90	0.90	0.90	10832
1	0.33	0.33	0.33	1611
micro avg	0.82	0.82	0.82	12443
macro avg	0.61	0.62	0.61	12443
weighted avg	0.83	0.82	0.83	12443

```
{'class_weight': 'balanced', 'criterion': 'entropy', 'max_depth': 20,
 'max_features': 'auto', 'min_samples_leaf': 1, 'min_samples_split': 2,
 'splitter': 'best'}
```

b. What is the size of tree (i.e. number of nodes)? Is the size different from the maximal tree or the tree in the previous step? Why?

In [21]:

```
print("Number of nodes: ", cv.best_estimator_.tree_.node_count)
```

Number of nodes: 13663

c. How many leaves are in the tree that is selected based on the validation dataset?

In []:

d. Which variable is used for the first split? What are the competing splits for this first split?

In [22]:

```
visualize_decision_tree(cv.best_estimator_, df.drop("IsBadBuy", axis=1).columns,
 "Tree_Struct_CV.png")
```

e. What are the 5 important variables in building the tree?

In [23]:

```
analyse_feature_importance(cv.best_estimator_, df.drop("IsBadBuy", axis=1).columns, 5)
```

```
MMRAcquisitionAuctionAveragePrice : 0.07556813524312803  
MMRAcquisitionRetailAveragePrice : 0.07453858047632252  
Veh0do : 0.07357030511029244  
MMRAcquisitionAuctionCleanPrice : 0.07355390254270788  
MMRAcquisitionRetailCleanPrice : 0.0672916389998756
```

f. Report if you see any evidence of model overfitting.

In []:

g. What are the parameters used? Explain your choices.

In []:

3. What is the significant difference do you see between these two decision tree models (steps 2.1 & 2.2)? How do they compare performance-wise? Explain why those changes may have happened.

In []:

4. From the better model, can you identify which cars could potential be “kicks”? Can you provide some descriptive summary of those cars?

In []:

In []:

Task 3. Predictive Modeling Using Regression

1. In preparation for regression, is any imputation of missing values needed for this data set? List the variables that needed this.

In [24]:

```
# We've already done this in the prep_data function
```

2. Apply transformation method(s) to the variable(s) that need it. List the variables that needed it

In [25]:

```

## Doing the log transformation

### Q: It's enoguh?
columns_to_transform = interval_cols

def logTransformation(df):

    df_log = df.copy()

    for col in columns_to_transform:
        df_log[col] = df_log[col].apply(lambda x: x+1)
        df_log[col] = df_log[col].apply(np.log)

    return df_log

df_log = logTransformation(df)
X_train_log, X_test_log, y_train_log, y_test_log = train_test_split(df_log.drop
(['IsBadBuy'], axis=1), df_log['IsBadBuy'], test_size=0.3, stratify=df_log['IsBa
dBuy'], random_state=rs)

if ResamplingMethod == 'ros':
    print("Using ROS Resmapling")
    ros = RandomOverSampler(random_state=rs)
    X_train_log, y_train_log = ros.fit_resample(X_train_log, y_train_log)
elif ResamplingMethod == 'rus':
    print("Using RUS Resmapling")
    rus = RandomUnderSampler(random_state=rs)
    X_train_log, y_train_log = rus.fit_resample(X_train_log, y_train_log)
else:
    print("No Resampling Method Used")

# Standardise
scaler_log = StandardScaler()
X_train_log = scaler_log.fit_transform(X_train_log, y_train_log)
X_test_log = scaler_log.transform(X_test_log)

```

Using ROS Resmapling

3. Build a regression model using the default regression method with all inputs. Once you done it, build another one and tune it using GridSearchCV. Answer the followings:

In [26]:

```

### Traing Logistic Regression
model = LogisticRegression(random_state=rs)
model.fit(X_train_log, y_train_log)

```

Out[26]:

```

LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
                    intercept_scaling=1, max_iter=100, multi_class='warn',
                    n_jobs=None, penalty='l2', random_state=101, solver='warn',
                    tol=0.0001, verbose=0, warm_start=False)

```

In [27]:

```

## GridSearch for Logistic Regression
params = {
    'C': [pow(10, x) for x in range(-4, 1)],
    'solver': ['newton-cg', 'lbfgs', 'liblinear', 'sag', 'saga'],
    'max_iter': [30, 50, 100],
    'warm_start': [True, False],
    'class_weight': ['balanced', None]
}

cv = GridSearchCV(param_grid=params, estimator=LogisticRegression(random_state=rs), cv=3, n_jobs=-1)
cv.fit(X_train_log, y_train_log)

```

Out[27]:

```

GridSearchCV(cv=3, error_score='raise-deprecating',
             estimator=LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
             intercept_scaling=1, max_iter=100, multi_class='warn',
             n_jobs=None, penalty='l2', random_state=101, solver='warn',
             tol=0.0001, verbose=0, warm_start=False),
             fit_params=None, iid='warn', n_jobs=-1,
             param_grid={'C': [0.0001, 0.001, 0.01, 0.1, 1], 'solver': ['newton-cg', 'lbfgs', 'liblinear', 'sag', 'saga'], 'max_iter': [30, 50, 100], 'warm_start': [True, False], 'class_weight': ['balanced', None]},
             pre_dispatch='2*n_jobs', refit=True, return_train_score='warn',
             scoring=None, verbose=0)

```

h. Name the regression function used.

In []:

i. How much was the difference in performance of two models build, default and optimal?

In [28]:

```
print("Train accuracy:", model.score(X_train_log, y_train_log))
print("Test accuracy:", model.score(X_test_log, y_test_log))
print("GridSearch Train accuracy:", cv.score(X_train_log, y_train_log))
print("GridSearch Test accuracy:", cv.score(X_test_log, y_test_log))
```

Train accuracy: 0.6967910418232897
 Test accuracy: 0.7531945672265531
 GridSearch Train accuracy: 0.6972460728841056
 GridSearch Test accuracy: 0.7545607972353934

j. Show the set parameters for the best model. What are the parameters used? Explain your decision. What are the optimal parameters?

In [29]:

```
print("The best model parameters: ", cv.best_params_)
```

The best model parameters: {'C': 0.1, 'class_weight': 'balanced', 'max_iter': 50, 'solver': 'lbfgs', 'warm_start': True}

k. Report which variables are included in the regression model.

In []:

l. Report the top-5 important variables (in the order) in the model.

In [30]:

```
def printLRTopImportant(model, top = 5):
    coef = model.coef_[0]
    indices = np.argsort(np.absolute(coef))
    indices = np.flip(indices, axis=0)
    indices = indices[:top]
    for i in indices:
        print(feature_names[i], ': ', coef[i])
```

In [31]:

```
printLRTopImportant(model, 5)
```

MMRAcquisitionAuctionAveragePrice : -1.2096083923720427
 MMRAcquisitionRetailAveragePrice : 1.172070760190353
 WheelTypeID_? : 0.7771546187078459
 WheelTypeID_1 : -0.6112154370497105
 WheelType_covers : -0.517388271054883

m. What is classification accuracy on training and test datasets?

In [32]:

```

y_pred = model.predict(X_test_log)
print("Classification Report: \n\n",classification_report(y_test_log, y_pred))

y_pred = cv.predict(X_test_log)
print("GridSearch Classification Report: \n\n",classification_report(y_test_log,
y_pred))
log_reg_model = cv.best_estimator_

```

Classification Report:

	precision	recall	f1-score	support
0	0.93	0.78	0.85	10832
1	0.29	0.60	0.39	1611
micro avg	0.75	0.75	0.75	12443
macro avg	0.61	0.69	0.62	12443
weighted avg	0.85	0.75	0.79	12443

GridSearch Classification Report:

	precision	recall	f1-score	support
0	0.93	0.78	0.85	10832
1	0.29	0.60	0.39	1611
micro avg	0.75	0.75	0.75	12443
macro avg	0.61	0.69	0.62	12443
weighted avg	0.85	0.75	0.79	12443

n. Report any sign of overfitting.

In [33]:

```
## The GridSearch Precision and Recall is weird
```

4. Build another regression model using the subset of inputs selected by RFE and selection by model method. Answer the followings:

In [34]:

```

rfe = RFECV(estimator = LogisticRegression(random_state=rs), cv=3)
rfe.fit(X_train_log, y_train_log)
X_train_rfe = rfe.transform(X_train_log)
X_test_rfe = rfe.transform(X_test_log)

selectmodel = SelectFromModel(dt_model, prefit=True)
X_train_sel_model = selectmodel.transform(X_train_log)
X_test_sel_model = selectmodel.transform(X_test_log)

```


a. Report which variables are included in the regression model.

In [35]:

```
print("Original feature set", X_train.shape[1])  
print("Number of RFE-selected features: ", rfe.n_features_)  
print("Number of selectFromModel features: ", X_train_sel_model.shape[1])
```

Original feature set 198

Number of RFE-selected features: 90

Number of selectFromModel features: 38

In [36]:

```
print("The RFE-selected features: \n\n", list(compress(feature_names, rfe.support_)))
print("\n\n")
print("The SelectFromModel features: \n\n", list(compress(feature_names, selectmodel.get_support())))
```

The RFE-selected features:

```
['VehOdo', 'MMRAcquisitionAuctionAveragePrice', 'MMRAcquisitionRetailAveragePrice', 'MMRAcquisitionRetailCleanPrice', 'VehBCost', 'WarrantyCost', 'Auction_adesa', 'VehYear_2001.0', 'VehYear_2002.0', 'VehYear_2003.0', 'VehYear_2004.0', 'VehYear_2006.0', 'VehYear_2007.0', 'VehYear_2008.0', 'VehYear_2009.0', 'Make_acura', 'Make_chrysler', 'Make_dodge', 'Make_ford', 'Make_honda', 'Make_infiniti', 'Make_jeep', 'Make_lexus', 'Make_nissan', 'Make_pontiac', 'Make_suzuki', 'Make_toyota', 'Make_volvo', 'WheelTypeID_0', 'WheelTypeID_1', 'WheelTypeID_3', 'WheelTypeID_?', 'WheelType_?', 'WheelType_alloy', 'WheelType_covers', 'WheelType_special', 'Nationality_other_asian', 'Nationality_top_line_asian', 'Size_large', 'Size_large_suv', 'Size_medium', 'Size_medium_suv', 'Size_van', 'TopThreeAmericanName_chrysler', 'TopThreeAmericanName_gm', 'MMRCurrentAuctionAveragePrice_5480', 'MMRCurrentAuctionAveragePrice_5569', 'MMRCurrentAuctionAveragePrice_6311', 'MMRCurrentAuctionAveragePrice_7269', 'MMRCurrentAuctionAveragePrice_7644', 'MMRCurrentAuctionAveragePrice_8186', 'MMRCurrentAuctionAveragePrice_LESS_FREQ', 'MMRCurrentAuctionCleanPrice_6461', 'MMRCurrentAuctionCleanPrice_6584', 'MMRCurrentAuctionCleanPrice_7450', 'MMRCurrentAuctionCleanPrice_7898', 'MMRCurrentAuctionCleanPrice_LESS_FREQ', 'MMRCurrentRetailAveragePrice_10834', 'MMRCurrentRetailAveragePrice_11674', 'MMRCurrentRetailAveragePrice_6418', 'MMRCurrentRetailAveragePrice_6515', 'MMRCurrentRetailAveragePrice_7316', 'MMRCurrentRetailAveragePrice_8756', 'MMRCurrentRetailAveragePrice_LESS_FREQ', 'MMRCurrentRetailCleanPrice_10103', 'MMRCurrentRetailCleanPrice_11413', 'MMRCurrentRetailCleanPrice_12387', 'MMRCurrentRetailCleanPrice_7478', 'MMRCurrentRetailCleanPrice_7611', 'MMRCurrentRetailCleanPrice_8546', 'MMRCurrentRetailCleanPrice_LESS_FREQ', 'MMRCurrentRetailRatio_#VALUE!', 'MMRCurrentRetailRatio_0.855997898', 'MMRCurrentRetailRatio_0.856073017', 'MMRCurrentRetailRatio_0.858250869', 'MMRCurrentRetailRatio_0.866673265', 'MMRCurrentRetailRatio_0.949268378', 'MMRCurrentRetailRatio_LESS_FREQ', 'PRIMEUNIT_?', 'PRIMEUNIT_no', 'VNST_id', 'VNST_ky', 'VNST_nc', 'VNST_ne', 'VNST_nh', 'VNST_ny', 'VNST_or', 'VNST_pa', 'VNST_tn', 'VNST_tx']
```

The SelectFromModel features:

```
['VehOdo', 'MMRAcquisitionAuctionAveragePrice', 'MMRAcquisitionAuctionCleanPrice', 'MMRAcquisitionRetailAveragePrice', 'MMRAcquisitionRetailCleanPrice', 'VehBCost', 'WarrantyCost', 'Auction_adesa', 'Auction_manheim', 'Auction_other', 'VehYear_2003.0', 'VehYear_2004.0', 'VehYear_2005.0', 'VehYear_2006.0', 'Make_chevrolet', 'Make_chrysler', 'Make_dodge', 'Make_ford', 'Color_black', 'Color_blue', 'Color_gold', 'Color_grey', 'Color_red', 'Color_silver', 'Color_white', 'WheelTypeID_?', 'WheelType_?', 'WheelType_alloy', 'WheelType_covers', 'Size_medium', 'TopThreeAmericanName_chrysler', 'TopThreeAmericanName_gm', 'VNST_az', 'VNST_ca', 'VNST_co', 'VNST_fl', 'VNST_nc', 'VNST_tx']
```

b. Report the top-5 important variables (in the order) in the model.

In [37]:

```
params = {
    'C': [pow(10, x) for x in range(-4, 1)],
    'solver' : ['newton-cg', "lbfgs", "liblinear", "sag", "saga"],
    'max_iter': [30, 50, 100],
    'warm_start': [True, False],
    'class_weight': ['balanced', None]
}
rfe_cv = GridSearchCV(param_grid=params, estimator=LogisticRegression(random_state=rs, verbose=True), cv=3, n_jobs=-1)
rfe_cv.fit(X_train_rfe, y_train_log)

selectModel_cv = GridSearchCV(param_grid=params, estimator=LogisticRegression(random_state=rs, verbose=True), cv=3, n_jobs=-1)
selectModel_cv.fit(X_train_sel_model, y_train_log)
```

```
[Parallel(n_jobs=1)]: Using backend SequentialBackend with 1 concurrent workers.
[Parallel(n_jobs=1)]: Done 1 out of 1 | elapsed: 2.6s finished
[Parallel(n_jobs=1)]: Using backend SequentialBackend with 1 concurrent workers.
[Parallel(n_jobs=1)]: Done 1 out of 1 | elapsed: 0.4s finished
```

Out[37]:

```
GridSearchCV(cv=3, error_score='raise-deprecating',
             estimator=LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
             intercept_scaling=1, max_iter=100, multi_class='warn',
             n_jobs=None, penalty='l2', random_state=101, solver='warn',
             tol=0.0001, verbose=True, warm_start=False),
             fit_params=None, iid='warn', n_jobs=-1,
             param_grid={'C': [0.0001, 0.001, 0.01, 0.1, 1], 'solver': ['newton-cg', 'lbfgs', 'liblinear', 'sag', 'saga'], 'max_iter': [30, 50, 100], 'warm_start': [True, False], 'class_weight': ['balanced', None]},
             pre_dispatch='2*n_jobs', refit=True, return_train_score='warn',
             scoring=None, verbose=0)
```

In [38]:

```
print("Top-5 important variables for RFE: \n")
printLRTopImportant(rfe_cv.best_estimator_, 5)
print("\n\n")
print("Top-5 important variables for selectModel \n")
printLRTopImportant(selectModel_cv.best_estimator_, 5)
```

Top-5 important variables for RFE:

```
MMRAcquisitionAuctionAveragePrice : -0.9620067587313315
MMRAcquisitionAuctionCleanPrice : 0.858021895106162
Make_infiniti : 0.5946791320209158
Make_honda : -0.4498818947004017
Make_isuzu : 0.44302155083592776
```

Top-5 important variables for selectModel

```
MMRAcquisitionAuctionAveragePrice : -1.307036409153068
MMRAcquisitionRetailCleanPrice : 1.0802608352180163
Make_chrysler : 0.6273796723439636
Make_dodge : 0.540067113709443
MMRAcquisitionAuctionCleanPrice : 0.3335244152330805
```

c. What are the parameters used? Explain your choices. What are the optimal parameters? Which regression function is being used?

In [39]:

```
print("Optimal Parameters for RFE", rfe_cv.best_params_)
print("Optimal Parameters for selectModel", selectModel_cv.best_params_)
```

```
Optimal Parameters for RFE {'C': 0.1, 'class_weight': 'balanced', 'max_iter': 30, 'solver': 'newton-cg', 'warm_start': True}
Optimal Parameters for selectModel {'C': 1, 'class_weight': 'balanced', 'max_iter': 50, 'solver': 'lbfgs', 'warm_start': True}
```

d. Report any sign of overfitting

In []:

e. What is classification accuracy on training and test datasets?

In [40]:

```
print("GridSearch Train accuracy:", cv.score(X_train_log, y_train_log))
print("GridSearch Test accuracy:", cv.score(X_test_log, y_test_log))
print("\n\nRFE:\n")
print("Train accuracy:", rfe_cv.score(X_train_rfe, y_train_log))
print("Test accuracy:", rfe_cv.score(X_test_rfe, y_test_log))
print("\n\nselectModel:\n")
print("Train accuracy:", selectModel_cv.score(X_train_sel_model, y_train_log))
print("Test accuracy:", selectModel_cv.score(X_test_sel_model, y_test_log))
```

GridSearch Train accuracy: 0.6972460728841056
GridSearch Test accuracy: 0.7545607972353934

RFE:

Train accuracy: 0.7000158271673327
Test accuracy: 0.7536767660532026

selectModel:

Train accuracy: 0.6827444308154949
Test accuracy: 0.7653299043638994

f. Did it improve/worsen the performance? Explain why those changes may have happened

In [41]:

```

y_pred = rfe_cv.predict(X_test_rfe)
print("REF classification report: \n",classification_report(y_test, y_pred))
print("\n\n")
y_pred = selectModel_cv.predict(X_test_sel_model)
print("selectModel classification report: \n",classification_report(y_test, y_pred))

```

REF classification report:

	precision	recall	f1-score	support
0	0.93	0.78	0.85	10832
1	0.29	0.60	0.39	1611
micro avg	0.75	0.75	0.75	12443
macro avg	0.61	0.69	0.62	12443
weighted avg	0.85	0.75	0.79	12443

selectModel classification report:

	precision	recall	f1-score	support
0	0.93	0.79	0.85	10832
1	0.29	0.57	0.39	1611
micro avg	0.77	0.77	0.77	12443
macro avg	0.61	0.68	0.62	12443
weighted avg	0.84	0.77	0.79	12443

Task4 - Predicting using neural network

1. Build a Neural Network model using the default setting. Answer the following:

In [42]:

```

model = MLPClassifier(random_state=rs)
model.fit(X_train_log, y_train_log)

```

Out[42]:

```

MLPClassifier(activation='relu', alpha=0.0001, batch_size='auto', beta_1=0.9,
              beta_2=0.999, early_stopping=False, epsilon=1e-08,
              hidden_layer_sizes=(100,), learning_rate='constant',
              learning_rate_init=0.001, max_iter=200, momentum=0.9,
              n_iter_no_change=10, nesterovs_momentum=True, power_t=0.5,
              random_state=101, shuffle=True, solver='adam', tol=0.0001,
              validation_fraction=0.1, verbose=False, warm_start=False)

```

a. What is the network architecture?

In [43]:

```
def printMLPArchitecture(model):

    print("Number of Layers: ",model.n_layers_ )
    print("The First layer is Input Layer, and the last layer is the output layer")
    for i, w in enumerate(model.coefs_):
        print("{} Layer with hidden size {}".format(i+1, w.shape[0]))
        if (i+1) == len(model.coefs_):
            print("{} Layer with hidden size {}".format(i+2, w.shape[1]))

    print("The activation function: ", model.activation)

printMLPArchitecture(model)
```

```
Number of Layers: 3
The First layer is Input Layer, and the last layer is the output layer
1 Layer with hidden size 198
2 Layer with hidden size 100
3 Layer with hidden size 1
The activation function: relu
```

b. How many iterations are needed to train this network?

In [44]:

```
print("Number of iterations it ran: ", model.n_iter_)
```

```
Number of iterations it ran: 200
```

c. Do you see any sign of over-fitting?

In [45]:

```
# fig = plt.figure(figsize=(10, 5))
# plt.ylabel('Accuracy', fontsize=15)
# plt.xlabel('Number of iterations', fontsize=15)
# plt.title('Validation Accuracy', fontsize=20, fontweight="bold")
# plt.plot(model.validation_scores_, label="Validation Accuracy")
```

d. Did the training process converge and resulted in the best model?


In [46]:

```
fig = plt.figure(figsize=(10, 5))
plt.ylabel('Loss',fontsize=15)
plt.xlabel('Number of iterations',fontsize=15)
plt.title('Training Loss',fontsize=20,fontweight = "bold")
plt.plot(model.loss_curve_, label="Training Loss")
```

The Loss curve is still decreasing

Out[46]:

[<matplotlib.lines.Line2D at 0x7fceb132edd8>]



e. What is classification accuracy on training and test datasets?

In [47]:

```
print("MLP Train accuracy:", model.score(X_train, y_train))
print("MLP Test accuracy:", model.score(X_test, y_test))
print("\n\n")
y_pred = model.predict(X_test)
print("MLP classification report: \n",classification_report(y_test, y_pred))
```

MLP Train accuracy: 0.4818383254856962

MLP Test accuracy: 0.7743309491280238

MLP classification report:

	precision	recall	f1-score	support
0	0.87	0.88	0.87	10832
1	0.09	0.09	0.09	1611
micro avg	0.77	0.77	0.77	12443
macro avg	0.48	0.48	0.48	12443
weighted avg	0.77	0.77	0.77	12443

2. Refine this network by tuning it with GridSearchCV.

In []:

In [48]:

```

# Default
# params = {'hidden_layer_sizes': [(3,), (5,), (7,), (9,)], 'alpha': [0.01, 0.001, 0.0001, 0.00001]}

params = [
    {
        'hidden_layer_sizes': [(128, 64, 32, 16)],
        'activation': ['relu'],
        'solver': ['adam'],
        'batch_size': [64],
        'shuffle': [True],
        'learning_rate_init': [0.001],
        'n_iter_no_change': [10],
        'max_iter': [200],
        'warm_start': [True],
        'early_stopping': [True],
        'alpha': [0.01, 0.001],
    },
]

cv = GridSearchCV(param_grid=params, estimator=MLPClassifier(random_state=rs, verbose=True), cv=3, n_jobs=-1)
# cv = GridSearchCV(param_grid=params, estimator=MLPClassifier(random_state=rs, early_stopping=True, max_iter = max_iter, n_iter_no_change = max_iter ), cv=3, n_jobs=-1)
cv.fit(X_train_log, y_train_log)

```

Iteration 1, loss = 0.54856218
Validation score: 0.731751
Iteration 2, loss = 0.47211536
Validation score: 0.771316
Iteration 3, loss = 0.40356883
Validation score: 0.799011
Iteration 4, loss = 0.34348921
Validation score: 0.834817
Iteration 5, loss = 0.29075779
Validation score: 0.851236
Iteration 6, loss = 0.24983120
Validation score: 0.860534
Iteration 7, loss = 0.22284442
Validation score: 0.876558
Iteration 8, loss = 0.20155821
Validation score: 0.888625
Iteration 9, loss = 0.18602661
Validation score: 0.882690
Iteration 10, loss = 0.16979396
Validation score: 0.898516
Iteration 11, loss = 0.15977150
Validation score: 0.896736
Iteration 12, loss = 0.15176313
Validation score: 0.895351
Iteration 13, loss = 0.14464423
Validation score: 0.905836
Iteration 14, loss = 0.13999550
Validation score: 0.907023
Iteration 15, loss = 0.12720528
Validation score: 0.906034
Iteration 16, loss = 0.12930383
Validation score: 0.907023
Iteration 17, loss = 0.12255444
Validation score: 0.910386
Iteration 18, loss = 0.12068214
Validation score: 0.911771
Iteration 19, loss = 0.11399424
Validation score: 0.908803
Iteration 20, loss = 0.11282158
Validation score: 0.915331
Iteration 21, loss = 0.10661677
Validation score: 0.922057
Iteration 22, loss = 0.10446701
Validation score: 0.916518
Iteration 23, loss = 0.10589664
Validation score: 0.909792
Iteration 24, loss = 0.10556269
Validation score: 0.918497
Iteration 25, loss = 0.09938828
Validation score: 0.918299
Iteration 26, loss = 0.09541842
Validation score: 0.912760
Iteration 27, loss = 0.10113910
Validation score: 0.919683
Iteration 28, loss = 0.09432959
Validation score: 0.909792
Iteration 29, loss = 0.09296471
Validation score: 0.923046
Iteration 30, loss = 0.09169923
Validation score: 0.924036
Iteration 31, loss = 0.08993748

Validation score: 0.919288
Iteration 32, loss = 0.08603412
Validation score: 0.921662
Iteration 33, loss = 0.08854639
Validation score: 0.922453
Iteration 34, loss = 0.08351747
Validation score: 0.922849
Iteration 35, loss = 0.08777318
Validation score: 0.917112
Iteration 36, loss = 0.08058952
Validation score: 0.928190
Iteration 37, loss = 0.08469264
Validation score: 0.923640
Iteration 38, loss = 0.08402573
Validation score: 0.931157
Iteration 39, loss = 0.07430225
Validation score: 0.926607
Iteration 40, loss = 0.07893320
Validation score: 0.917903
Iteration 41, loss = 0.07689212
Validation score: 0.925223
Iteration 42, loss = 0.07355570
Validation score: 0.922057
Iteration 43, loss = 0.08132868
Validation score: 0.927399
Iteration 44, loss = 0.07889856
Validation score: 0.923838
Iteration 45, loss = 0.07490265
Validation score: 0.927596
Iteration 46, loss = 0.07099558
Validation score: 0.928388
Iteration 47, loss = 0.06990805
Validation score: 0.931157
Iteration 48, loss = 0.07486228
Validation score: 0.929179
Iteration 49, loss = 0.07158638
Validation score: 0.931355
Iteration 50, loss = 0.06874335
Validation score: 0.923046
Iteration 51, loss = 0.07661840
Validation score: 0.924431
Iteration 52, loss = 0.06855482
Validation score: 0.918101
Iteration 53, loss = 0.07291045
Validation score: 0.927596
Iteration 54, loss = 0.06441686
Validation score: 0.918892
Iteration 55, loss = 0.06992323
Validation score: 0.934520
Iteration 56, loss = 0.07443742
Validation score: 0.929970
Iteration 57, loss = 0.06685328
Validation score: 0.930366
Iteration 58, loss = 0.06478549
Validation score: 0.928586
Iteration 59, loss = 0.06425201
Validation score: 0.920079
Iteration 60, loss = 0.06606451
Validation score: 0.926805
Iteration 61, loss = 0.06558700
Validation score: 0.917507

```

Iteration 62, loss = 0.06666644
Validation score: 0.928586
Iteration 63, loss = 0.06158633
Validation score: 0.933927
Iteration 64, loss = 0.06551972
Validation score: 0.928190
Iteration 65, loss = 0.06410788
Validation score: 0.933927
Iteration 66, loss = 0.06246794
Validation score: 0.931355
Validation score did not improve more than tol=0.000100 for 10 consecutive epochs. Stopping.

```

Out[48]:

```

GridSearchCV(cv=3, error_score='raise-deprecating',
             estimator=MLPClassifier(activation='relu', alpha=0.0001, batch_size='auto', beta_1=0.9,
                                     beta_2=0.999, early_stopping=False, epsilon=1e-08,
                                     hidden_layer_sizes=(100,), learning_rate='constant',
                                     learning_rate_init=0.001, max_iter=200, momentum=0.9,
                                     n_iter_no_change=10, nesterovs_momentum=True, power_t=0.5,
                                     random_state=101, shuffle=True, solver='adam', tol=0.0001,
                                     validation_fraction=0.1, verbose=True, warm_start=False),
             fit_params=None, iid='warn', n_jobs=-1,
             param_grid=[{'hidden_layer_sizes': [(128, 64, 32, 16)], 'activation': ['relu'], 'solver': ['adam'], 'batch_size': [64], 'shuffle': [True], 'learning_rate_init': [0.001], 'n_iter_no_change': [10], 'max_iter': [200], 'warm_start': [True], 'early_stopping': [True], 'alpha': [0.01, 0.001]}],
             pre_dispatch='2*n_jobs', refit=True, return_train_score='warn',
             scoring=None, verbose=0)

```

a. What is the network architecture?

In [53]:

```
printMLPArchitecture(cv.best_estimator_)
```

```

Number of Layers: 6
The First layer is Input Layer, and the last layer is the output layer
1 Layer with hidden size 198
2 Layer with hidden size 128
3 Layer with hidden size 64
4 Layer with hidden size 32
5 Layer with hidden size 16
6 Layer with hidden size 1
The activation function: relu

```

b. How many iterations are needed to train this network?

In [54]:

```
print("Number of iterations it ran: ",cv.best_estimator_.n_iter_)
```

Number of iterations it ran: 66


c. Sign of overfitting?

In [55]:

```
fig = plt.figure(figsize=(10, 5))  
plt.ylabel('Accuracy',fontsize=15)  
plt.xlabel('Number of iterations',fontsize=15)  
plt.title('Validation Accuracy',fontsize=20,fontweight = "bold")  
plt.plot(cv.best_estimator_.validation_scores_, label="Validation Accuracy")
```

Out[55]:

[<matplotlib.lines.Line2D at 0x7fceb73ae240>]




d. Did the training process converge and resulted in the best model?

In [56]:

```
fig = plt.figure(figsize=(10, 5))
plt.ylabel('Loss',fontsize=15)
plt.xlabel('Number of iterations',fontsize=15)
plt.title('Training Loss',fontsize=20,fontweight = "bold")
plt.plot(cv.best_estimator_.loss_curve_, label="Training Loss")
```

Out[56]:

[<matplotlib.lines.Line2D at 0x7fcee981feb8>]



e. What is classification accuracy on training and test datasets? Is there any improvement in the outcome?

In [57]:

```
print("GridSearch NN Train accuracy:", cv.score(X_train_log, y_train_log))
print("GridSearch NN Test accuracy:", cv.score(X_test_log, y_test_log))

print("\n\n")
y_pred = cv.predict(X_test_log)
print("GridSearch NN Classification Report: \n", classification_report(y_test_log
, y_pred))

print("Best Parameters of NN: ", cv.best_params_)
nn_model = cv.best_estimator_
```

GridSearch NN Train accuracy: 0.975843785858426

GridSearch NN Test accuracy: 0.8350880012858636

GridSearch NN Classification Report:

	precision	recall	f1-score	support
0	0.90	0.91	0.91	10832
1	0.36	0.35	0.36	1611
micro avg	0.84	0.84	0.84	12443
macro avg	0.63	0.63	0.63	12443
weighted avg	0.83	0.84	0.83	12443

Best Parameters of NN: {'activation': 'relu', 'alpha': 0.001, 'batch_size': 64, 'early_stopping': True, 'hidden_layer_sizes': (128, 64, 32, 16), 'learning_rate_init': 0.001, 'max_iter': 200, 'n_iter_no_change': 10, 'shuffle': True, 'solver': 'adam', 'warm_start': True}

3. Would feature selection help here? Build another Neural Network model with inputs selected from RFE with regression (use the best model generated in Task 3) and selection with decision tree (use the best model from Task 2).

In [58]:

```
params = [
    {
        'hidden_layer_sizes': [(128, 64, 32, 16)],
        'activation': ['relu'],
        'solver' : ['adam'],
        'batch_size': [64],
        'shuffle': [True],
        'learning_rate_init': [0.001],
        'n_iter_no_change': [10],
        'max_iter':[200],
        'warm_start': [True],
        'early_stopping': [True],
        'alpha': [0.01, 0.001],
    },
]

rfe_cv = GridSearchCV(param_grid=params, estimator=MLPClassifier(random_state=rs
, early_stopping=True, verbose=True), cv=3, n_jobs=-1)
rfe_cv.fit(X_train_rfe, y_train_log)
modelSelect_cv = GridSearchCV(param_grid=params, estimator=MLPClassifier(random_
state=rs, early_stopping=True, verbose=True), cv=3, n_jobs=-1)
modelSelect_cv.fit(X_train_sel_model, y_train_log)
```

Iteration 1, loss = 0.55920437
Validation score: 0.701088
Iteration 2, loss = 0.52882666
Validation score: 0.718497
Iteration 3, loss = 0.50945268
Validation score: 0.724431
Iteration 4, loss = 0.49081665
Validation score: 0.733729
Iteration 5, loss = 0.47221045
Validation score: 0.748764
Iteration 6, loss = 0.45545179
Validation score: 0.756874
Iteration 7, loss = 0.43961116
Validation score: 0.762413
Iteration 8, loss = 0.42345590
Validation score: 0.776855
Iteration 9, loss = 0.40902212
Validation score: 0.774085
Iteration 10, loss = 0.39399478
Validation score: 0.782987
Iteration 11, loss = 0.38196869
Validation score: 0.791691
Iteration 12, loss = 0.37132417
Validation score: 0.802770
Iteration 13, loss = 0.35926943
Validation score: 0.802374
Iteration 14, loss = 0.35237085
Validation score: 0.807913
Iteration 15, loss = 0.34211636
Validation score: 0.815430
Iteration 16, loss = 0.33492760
Validation score: 0.819980
Iteration 17, loss = 0.32979855
Validation score: 0.813452
Iteration 18, loss = 0.32135280
Validation score: 0.821761
Iteration 19, loss = 0.31107574
Validation score: 0.820772
Iteration 20, loss = 0.30655272
Validation score: 0.828487
Iteration 21, loss = 0.30055642
Validation score: 0.834619
Iteration 22, loss = 0.29370566
Validation score: 0.839763
Iteration 23, loss = 0.29085304
Validation score: 0.840752
Iteration 24, loss = 0.28560413
Validation score: 0.839565
Iteration 25, loss = 0.27783852
Validation score: 0.848467
Iteration 26, loss = 0.27460390
Validation score: 0.843126
Iteration 27, loss = 0.26903037
Validation score: 0.849654
Iteration 28, loss = 0.26691435
Validation score: 0.846093
Iteration 29, loss = 0.26184186
Validation score: 0.855391
Iteration 30, loss = 0.25517697
Validation score: 0.857765
Iteration 31, loss = 0.25043793

Validation score: 0.860732
Iteration 32, loss = 0.24956826
Validation score: 0.867458
Iteration 33, loss = 0.24653509
Validation score: 0.864095
Iteration 34, loss = 0.24232718
Validation score: 0.864095
Iteration 35, loss = 0.23811419
Validation score: 0.856578
Iteration 36, loss = 0.23480829
Validation score: 0.867458
Iteration 37, loss = 0.23060738
Validation score: 0.871217
Iteration 38, loss = 0.23409495
Validation score: 0.864688
Iteration 39, loss = 0.22532895
Validation score: 0.878734
Iteration 40, loss = 0.22178624
Validation score: 0.876558
Iteration 41, loss = 0.22148049
Validation score: 0.876558
Iteration 42, loss = 0.21450477
Validation score: 0.865084
Iteration 43, loss = 0.21584733
Validation score: 0.874777
Iteration 44, loss = 0.20695801
Validation score: 0.869436
Iteration 45, loss = 0.21020165
Validation score: 0.871414
Iteration 46, loss = 0.20530024
Validation score: 0.883284
Iteration 47, loss = 0.20353714
Validation score: 0.883086
Iteration 48, loss = 0.20482862
Validation score: 0.875964
Iteration 49, loss = 0.19887822
Validation score: 0.877151
Iteration 50, loss = 0.19629105
Validation score: 0.877745
Iteration 51, loss = 0.19356330
Validation score: 0.884866
Iteration 52, loss = 0.19437609
Validation score: 0.885460
Iteration 53, loss = 0.18961563
Validation score: 0.879130
Iteration 54, loss = 0.19083436
Validation score: 0.884075
Iteration 55, loss = 0.18806997
Validation score: 0.887240
Iteration 56, loss = 0.18442706
Validation score: 0.889416
Iteration 57, loss = 0.18377619
Validation score: 0.883482
Iteration 58, loss = 0.18678629
Validation score: 0.889021
Iteration 59, loss = 0.18137343
Validation score: 0.887438
Iteration 60, loss = 0.18123354
Validation score: 0.888823
Iteration 61, loss = 0.17177968
Validation score: 0.885064

Iteration 62, loss = 0.17651974
Validation score: 0.884669
Iteration 63, loss = 0.17674598
Validation score: 0.886053
Iteration 64, loss = 0.17642252
Validation score: 0.881701
Iteration 65, loss = 0.17523190
Validation score: 0.890406
Iteration 66, loss = 0.16765737
Validation score: 0.882493
Iteration 67, loss = 0.17072475
Validation score: 0.882097
Iteration 68, loss = 0.17133876
Validation score: 0.885064
Iteration 69, loss = 0.16613050
Validation score: 0.889614
Iteration 70, loss = 0.16533149
Validation score: 0.897527
Iteration 71, loss = 0.16759448
Validation score: 0.899703
Iteration 72, loss = 0.16446430
Validation score: 0.897132
Iteration 73, loss = 0.16026355
Validation score: 0.900890
Iteration 74, loss = 0.16050879
Validation score: 0.895153
Iteration 75, loss = 0.15964436
Validation score: 0.890406
Iteration 76, loss = 0.16012514
Validation score: 0.898912
Iteration 77, loss = 0.16244396
Validation score: 0.895549
Iteration 78, loss = 0.15592813
Validation score: 0.895945
Iteration 79, loss = 0.15615357
Validation score: 0.892384
Iteration 80, loss = 0.15800574
Validation score: 0.895747
Iteration 81, loss = 0.15256501
Validation score: 0.904649
Iteration 82, loss = 0.15300054
Validation score: 0.902671
Iteration 83, loss = 0.15487216
Validation score: 0.896934
Iteration 84, loss = 0.14603086
Validation score: 0.908012
Iteration 85, loss = 0.15403320
Validation score: 0.899703
Iteration 86, loss = 0.15071394
Validation score: 0.894758
Iteration 87, loss = 0.14822653
Validation score: 0.898121
Iteration 88, loss = 0.14979554
Validation score: 0.893769
Iteration 89, loss = 0.15398114
Validation score: 0.899308
Iteration 90, loss = 0.14342492
Validation score: 0.904451
Iteration 91, loss = 0.13827129
Validation score: 0.908012
Iteration 92, loss = 0.13878313

Validation score: 0.901682
Iteration 93, loss = 0.14959304
Validation score: 0.901484
Iteration 94, loss = 0.14540613
Validation score: 0.901286
Iteration 95, loss = 0.13803944
Validation score: 0.899703
Validation score did not improve more than tol=0.000100 for 10 consecutive epochs. Stopping.
Iteration 1, loss = 0.56758633
Validation score: 0.708803
Iteration 2, loss = 0.52061289
Validation score: 0.742235
Iteration 3, loss = 0.47988162
Validation score: 0.765381
Iteration 4, loss = 0.43770721
Validation score: 0.789515
Iteration 5, loss = 0.40139163
Validation score: 0.802967
Iteration 6, loss = 0.36874125
Validation score: 0.822354
Iteration 7, loss = 0.33961134
Validation score: 0.837389
Iteration 8, loss = 0.31134099
Validation score: 0.835015
Iteration 9, loss = 0.29042468
Validation score: 0.857567
Iteration 10, loss = 0.27402650
Validation score: 0.854995
Iteration 11, loss = 0.25834679
Validation score: 0.853017
Iteration 12, loss = 0.24560960
Validation score: 0.865282
Iteration 13, loss = 0.23563915
Validation score: 0.868249
Iteration 14, loss = 0.22484281
Validation score: 0.871019
Iteration 15, loss = 0.21206894
Validation score: 0.876558
Iteration 16, loss = 0.20650296
Validation score: 0.881108
Iteration 17, loss = 0.19491230
Validation score: 0.889219
Iteration 18, loss = 0.19157948
Validation score: 0.891395
Iteration 19, loss = 0.18691339
Validation score: 0.898714
Iteration 20, loss = 0.17542886
Validation score: 0.890801
Iteration 21, loss = 0.16974231
Validation score: 0.895153
Iteration 22, loss = 0.16926681
Validation score: 0.885262
Iteration 23, loss = 0.16494753
Validation score: 0.896340
Iteration 24, loss = 0.15958589
Validation score: 0.893373
Iteration 25, loss = 0.15443237
Validation score: 0.896340
Iteration 26, loss = 0.15332352
Validation score: 0.905242

Iteration 27, loss = 0.14638877
Validation score: 0.902473
Iteration 28, loss = 0.14652707
Validation score: 0.896736
Iteration 29, loss = 0.13801102
Validation score: 0.905045
Iteration 30, loss = 0.13734004
Validation score: 0.904055
Iteration 31, loss = 0.13424590
Validation score: 0.909001
Iteration 32, loss = 0.13010395
Validation score: 0.909001
Iteration 33, loss = 0.13167964
Validation score: 0.904055
Iteration 34, loss = 0.12470479
Validation score: 0.908803
Iteration 35, loss = 0.12287827
Validation score: 0.916518
Iteration 36, loss = 0.12482632
Validation score: 0.909792
Iteration 37, loss = 0.12250820
Validation score: 0.910386
Iteration 38, loss = 0.11735599
Validation score: 0.919090
Iteration 39, loss = 0.11516125
Validation score: 0.916518
Iteration 40, loss = 0.11587339
Validation score: 0.915727
Iteration 41, loss = 0.11253537
Validation score: 0.913155
Iteration 42, loss = 0.11271207
Validation score: 0.912957
Iteration 43, loss = 0.10741468
Validation score: 0.918497
Iteration 44, loss = 0.10654170
Validation score: 0.907616
Iteration 45, loss = 0.11138189
Validation score: 0.912166
Iteration 46, loss = 0.10459638
Validation score: 0.914342
Iteration 47, loss = 0.10066200
Validation score: 0.919683
Iteration 48, loss = 0.10282604
Validation score: 0.919881
Iteration 49, loss = 0.10288856
Validation score: 0.919090
Iteration 50, loss = 0.09895830
Validation score: 0.916518
Iteration 51, loss = 0.10345620
Validation score: 0.921464
Iteration 52, loss = 0.09593454
Validation score: 0.914936
Iteration 53, loss = 0.09924745
Validation score: 0.912562
Iteration 54, loss = 0.09504412
Validation score: 0.924629
Iteration 55, loss = 0.09313935
Validation score: 0.920673
Iteration 56, loss = 0.09394586
Validation score: 0.917507
Iteration 57, loss = 0.09315643

Validation score: 0.913947
Iteration 58, loss = 0.09688742
Validation score: 0.920870
Iteration 59, loss = 0.09098033
Validation score: 0.926805
Iteration 60, loss = 0.08448669
Validation score: 0.913947
Iteration 61, loss = 0.09126251
Validation score: 0.920870
Iteration 62, loss = 0.08467944
Validation score: 0.917903
Iteration 63, loss = 0.08991536
Validation score: 0.923244
Iteration 64, loss = 0.08036737
Validation score: 0.916320
Iteration 65, loss = 0.09312787
Validation score: 0.929773
Iteration 66, loss = 0.08405142
Validation score: 0.924036
Iteration 67, loss = 0.08724836
Validation score: 0.927201
Iteration 68, loss = 0.08187973
Validation score: 0.928586
Iteration 69, loss = 0.08966199
Validation score: 0.927992
Iteration 70, loss = 0.08373264
Validation score: 0.920277
Iteration 71, loss = 0.08257443
Validation score: 0.924431
Iteration 72, loss = 0.07883727
Validation score: 0.924036
Iteration 73, loss = 0.07999067
Validation score: 0.922651
Iteration 74, loss = 0.08826804
Validation score: 0.924233
Iteration 75, loss = 0.07661485
Validation score: 0.912166
Iteration 76, loss = 0.08053650
Validation score: 0.918497
Validation score did not improve more than tol=0.000100 for 10 consecutive epochs. Stopping.

Out[58]:

```
GridSearchCV(cv=3, error_score='raise-deprecating',
             estimator=MLPClassifier(activation='relu', alpha=0.0001, batch_size='auto', beta_1=0.9,
                                     beta_2=0.999, early_stopping=True, epsilon=1e-08,
                                     hidden_layer_sizes=(100,), learning_rate='constant',
                                     learning_rate_init=0.001, max_iter=200, momentum=0.9,
                                     n_iter_no_change=10, nesterovs_momentum=True, power_t=0.5,
                                     random_state=101, shuffle=True, solver='adam', tol=0.0001,
                                     validation_fraction=0.1, verbose=True, warm_start=False),
             fit_params=None, iid='warn', n_jobs=-1,
             param_grid=[{'hidden_layer_sizes': [(128, 64, 32, 16)], 'activation': ['relu'], 'solver': ['adam'], 'batch_size': [64], 'shuffle': [True], 'learning_rate_init': [0.001], 'n_iter_no_change': [10], 'max_iter': [200], 'warm_start': [True], 'early_stopping': [True], 'alpha': [0.01, 0.001]}],
             pre_dispatch='2*n_jobs', refit=True, return_train_score='warn',
             scoring=None, verbose=0)
```

a. Did feature selection help here? Any change in the network architecture? What inputs are being used as the network input?

In [59]:

```
print("Best Parameters of NN: ", cv.best_params_)
print("Best Parameters of RFE NN: ", rfe_cv.best_params_)
print("Best Parameters of modelSelect NN: ", modelSelect_cv.best_params_)
print("\n\n")

print("GridSearch:")
printMLPArchitecture(cv.best_estimator_)
print("\n")
print("RFE:")
printMLPArchitecture(rfe_cv.best_estimator_)
print("\n")
print("modelSelect:")
printMLPArchitecture(modelSelect_cv.best_estimator_)
print("\n")
```

```
Best Parameters of NN: {'activation': 'relu', 'alpha': 0.001, 'batch_size': 64, 'early_stopping': True, 'hidden_layer_sizes': (128, 64, 32, 16), 'learning_rate_init': 0.001, 'max_iter': 200, 'n_iter_no_change': 10, 'shuffle': True, 'solver': 'adam', 'warm_start': True}
Best Parameters of RFE NN: {'activation': 'relu', 'alpha': 0.001, 'batch_size': 64, 'early_stopping': True, 'hidden_layer_sizes': (128, 64, 32, 16), 'learning_rate_init': 0.001, 'max_iter': 200, 'n_iter_no_change': 10, 'shuffle': True, 'solver': 'adam', 'warm_start': True}
Best Parameters of modelSelect NN: {'activation': 'relu', 'alpha': 0.001, 'batch_size': 64, 'early_stopping': True, 'hidden_layer_sizes': (128, 64, 32, 16), 'learning_rate_init': 0.001, 'max_iter': 200, 'n_iter_no_change': 10, 'shuffle': True, 'solver': 'adam', 'warm_start': True}
```

GridSearch:

Number of Layers: 6

The First layer is Input Layer, and the last layer is the output layer

- 1 Layer with hidden size 198
- 2 Layer with hidden size 128
- 3 Layer with hidden size 64
- 4 Layer with hidden size 32
- 5 Layer with hidden size 16
- 6 Layer with hidden size 1

The activation function: relu

RFE:

Number of Layers: 6

The First layer is Input Layer, and the last layer is the output layer

- 1 Layer with hidden size 90
- 2 Layer with hidden size 128
- 3 Layer with hidden size 64
- 4 Layer with hidden size 32
- 5 Layer with hidden size 16
- 6 Layer with hidden size 1

The activation function: relu

modelSelect:

Number of Layers: 6

The First layer is Input Layer, and the last layer is the output layer

- 1 Layer with hidden size 38
- 2 Layer with hidden size 128
- 3 Layer with hidden size 64
- 4 Layer with hidden size 32
- 5 Layer with hidden size 16
- 6 Layer with hidden size 1

The activation function: relu

b. What is classification accuracy on training and test datasets? Is there any improvement in the outcome?

In [60]:

```
print("GridSearch NN Train accuracy:", cv.score(X_train_log, y_train_log))
print("GridSearch NN Test accuracy:", cv.score(X_test_log, y_test_log))
print("RFE NN Train accuracy:", rfe_cv.score(X_train_rfe, y_train_log))
print("RFE NNTest accuracy:", rfe_cv.score(X_test_rfe, y_test_log))
print("modelSelect NN Train accuracy:", modelSelect_cv.score(X_train_sel_model,
y_train_log))
print("modelSelect NN Test accuracmodelSelect_cvy:", modelSelect_cv.score(X_test
_sel_model, y_test_log))
```

```
GridSearch NN Train accuracy: 0.975843785858426
GridSearch NN Test accuracy: 0.8350880012858636
RFE NN Train accuracy: 0.9499267993510861
RFE NNTest accuracy: 0.8128264887888773
modelSelect NN Train accuracy: 0.974102797451826
modelSelect NN Test accuracmodelSelect_cvy: 0.8151571164510166
```

c. How many iterations are now needed to train this network?

In [61]:

```
print("Number of iterations GS ran: ",cv.best_estimator_.n_iter_)
print("Number of iterations rfe ran: ",rfe_cv.best_estimator_.n_iter_)
print("Number of iterations modelSelect ran: ",modelSelect_cv.best_estimator_.n_
iter_)
```

```
Number of iterations GS ran: 66
Number of iterations rfe ran: 95
Number of iterations modelSelect ran: 76
```

d. Do you see any sign of over-fitting?

In []:


e. Did the training process converge and resulted in the best model?

In [62]:

```
fig = plt.figure(figsize=(10, 5))
plt.ylabel('Loss',fontsize=15)
plt.xlabel('Number of iterations',fontsize=15)
plt.title('Training Loss',fontsize=20,fontweight = "bold")
plt.plot(cv.best_estimator_.loss_curve_, label="GridSearch Loss")
plt.plot(rfe_cv.best_estimator_.loss_curve_, label="RFE Loss")
plt.plot(modelSelect_cv.best_estimator_.loss_curve_, label="modelSelect Loss")
plt.legend(loc='upper right')
```

Out[62]:

<matplotlib.legend.Legend at 0x7fcea6f45470>



4. Using the comparison methods, which of the models (i.e one with selected variables and another with all variables) appears to be better? From the better model, can you identify cars those could potential be “kicks”? Can you provide some descriptive summary of those cars? Is it easy to comprehend the performance of the best neural network model for decision making?

In [63]:

```
print("GridSearch Classification Report: ")
y_pred = cv.predict(X_test_log)
print(classification_report(y_test_log, y_pred))
print("\n\nRFE Classification Report: ")
y_pred = rfe_cv.predict(X_test_rfe)
print(classification_report(y_test_log, y_pred))
print("\n\nmodelSelect Classification Report: ")
y_pred = modelSelect_cv.predict(X_test_sel_model)
print(classification_report(y_test_log, y_pred))
```

GridSearch Classification Report:

	precision	recall	f1-score	support
0	0.90	0.91	0.91	10832
1	0.36	0.35	0.36	1611
micro avg	0.84	0.84	0.84	12443
macro avg	0.63	0.63	0.63	12443
weighted avg	0.83	0.84	0.83	12443

RFE Classification Report:

	precision	recall	f1-score	support
0	0.91	0.88	0.89	10832
1	0.32	0.39	0.35	1611
micro avg	0.81	0.81	0.81	12443
macro avg	0.61	0.63	0.62	12443
weighted avg	0.83	0.81	0.82	12443

modelSelect Classification Report:

	precision	recall	f1-score	support
0	0.90	0.89	0.89	10832
1	0.31	0.34	0.32	1611
micro avg	0.82	0.82	0.82	12443
macro avg	0.60	0.61	0.61	12443
weighted avg	0.82	0.82	0.82	12443

Task 5. Generating an Ensemble Model and Comparing Models

1. Generate an ensemble model to include the best regression model, best decision tree model, and best neural network model.

In [64]:

```
voting = VotingClassifier(estimators=[('dt', dt_model), ('lr', log_reg_model), ('nn', nn_model)], voting='soft')
voting.fit(X_train_log, y_train_log)

y_pred_dt = dt_model.predict(X_test_log)
y_pred_log_reg = log_reg_model.predict(X_test_log)
y_pred_nn = nn_model.predict(X_test_log)
y_pred_ensemble = voting.predict(X_test_log)
```

Iteration 1, loss = 0.54856218
Validation score: 0.731751
Iteration 2, loss = 0.47211536
Validation score: 0.771316
Iteration 3, loss = 0.40356883
Validation score: 0.799011
Iteration 4, loss = 0.34348921
Validation score: 0.834817
Iteration 5, loss = 0.29075779
Validation score: 0.851236
Iteration 6, loss = 0.24983120
Validation score: 0.860534
Iteration 7, loss = 0.22284442
Validation score: 0.876558
Iteration 8, loss = 0.20155821
Validation score: 0.888625
Iteration 9, loss = 0.18602661
Validation score: 0.882690
Iteration 10, loss = 0.16979396
Validation score: 0.898516
Iteration 11, loss = 0.15977150
Validation score: 0.896736
Iteration 12, loss = 0.15176313
Validation score: 0.895351
Iteration 13, loss = 0.14464423
Validation score: 0.905836
Iteration 14, loss = 0.13999550
Validation score: 0.907023
Iteration 15, loss = 0.12720528
Validation score: 0.906034
Iteration 16, loss = 0.12930383
Validation score: 0.907023
Iteration 17, loss = 0.12255444
Validation score: 0.910386
Iteration 18, loss = 0.12068214
Validation score: 0.911771
Iteration 19, loss = 0.11399424
Validation score: 0.908803
Iteration 20, loss = 0.11282158
Validation score: 0.915331
Iteration 21, loss = 0.10661677
Validation score: 0.922057
Iteration 22, loss = 0.10446701
Validation score: 0.916518
Iteration 23, loss = 0.10589664
Validation score: 0.909792
Iteration 24, loss = 0.10556269
Validation score: 0.918497
Iteration 25, loss = 0.09938828
Validation score: 0.918299
Iteration 26, loss = 0.09541842
Validation score: 0.912760
Iteration 27, loss = 0.10113910
Validation score: 0.919683
Iteration 28, loss = 0.09432959
Validation score: 0.909792
Iteration 29, loss = 0.09296471
Validation score: 0.923046
Iteration 30, loss = 0.09169923
Validation score: 0.924036
Iteration 31, loss = 0.08993748

Validation score: 0.919288
Iteration 32, loss = 0.08603412
Validation score: 0.921662
Iteration 33, loss = 0.08854639
Validation score: 0.922453
Iteration 34, loss = 0.08351747
Validation score: 0.922849
Iteration 35, loss = 0.08777318
Validation score: 0.917112
Iteration 36, loss = 0.08058952
Validation score: 0.928190
Iteration 37, loss = 0.08469264
Validation score: 0.923640
Iteration 38, loss = 0.08402573
Validation score: 0.931157
Iteration 39, loss = 0.07430225
Validation score: 0.926607
Iteration 40, loss = 0.07893320
Validation score: 0.917903
Iteration 41, loss = 0.07689212
Validation score: 0.925223
Iteration 42, loss = 0.07355570
Validation score: 0.922057
Iteration 43, loss = 0.08132868
Validation score: 0.927399
Iteration 44, loss = 0.07889856
Validation score: 0.923838
Iteration 45, loss = 0.07490265
Validation score: 0.927596
Iteration 46, loss = 0.07099558
Validation score: 0.928388
Iteration 47, loss = 0.06990805
Validation score: 0.931157
Iteration 48, loss = 0.07486228
Validation score: 0.929179
Iteration 49, loss = 0.07158638
Validation score: 0.931355
Iteration 50, loss = 0.06874335
Validation score: 0.923046
Iteration 51, loss = 0.07661840
Validation score: 0.924431
Iteration 52, loss = 0.06855482
Validation score: 0.918101
Iteration 53, loss = 0.07291045
Validation score: 0.927596
Iteration 54, loss = 0.06441686
Validation score: 0.918892
Iteration 55, loss = 0.06992323
Validation score: 0.934520
Iteration 56, loss = 0.07443742
Validation score: 0.929970
Iteration 57, loss = 0.06685328
Validation score: 0.930366
Iteration 58, loss = 0.06478549
Validation score: 0.928586
Iteration 59, loss = 0.06425201
Validation score: 0.920079
Iteration 60, loss = 0.06606451
Validation score: 0.926805
Iteration 61, loss = 0.06558700
Validation score: 0.917507


```
Iteration 62, loss = 0.06666644
Validation score: 0.928586
Iteration 63, loss = 0.06158633
Validation score: 0.933927
Iteration 64, loss = 0.06551972
Validation score: 0.928190
Iteration 65, loss = 0.06410788
Validation score: 0.933927
Iteration 66, loss = 0.06246794
Validation score: 0.931355
Validation score did not improve more than tol=0.000100 for 10 consecutive epochs. Stopping.
```

a. Does the Ensemble model outperform the underlying models? Resonate your answer.

In [65]:

```
print("Report for DT: \n",classification_report(y_test_log, y_pred_dt))
print("\nReport for Logistic Regression: \n",classification_report(y_test_log, y
_pred_log_reg))
print("\nReport for NN: \n",classification_report(y_test_log, y_pred_nn))
print("\nReport for Ensemble: \n",classification_report(y_test_log, y_pred_ensem
ble))
```

Report for DT:

	precision	recall	f1-score	support
0	0.88	1.00	0.93	10832
1	0.80	0.05	0.09	1611
micro avg	0.88	0.88	0.88	12443
macro avg	0.84	0.52	0.51	12443
weighted avg	0.87	0.88	0.82	12443

Report for Logistic Regression:

	precision	recall	f1-score	support
0	0.93	0.78	0.85	10832
1	0.29	0.60	0.39	1611
micro avg	0.75	0.75	0.75	12443
macro avg	0.61	0.69	0.62	12443
weighted avg	0.85	0.75	0.79	12443

Report for NN:

	precision	recall	f1-score	support
0	0.90	0.91	0.91	10832
1	0.36	0.35	0.36	1611
micro avg	0.84	0.84	0.84	12443
macro avg	0.63	0.63	0.63	12443
weighted avg	0.83	0.84	0.83	12443

Report for Ensemble:

	precision	recall	f1-score	support
0	0.91	0.93	0.92	10832
1	0.43	0.37	0.40	1611
micro avg	0.85	0.85	0.85	12443
macro avg	0.67	0.65	0.66	12443
weighted avg	0.85	0.85	0.85	12443

2. Use the comparison methods (or the comparison node) to compare the best decision tree model, the best regression model, the best neural network model and the ensemble model.

a. Discuss the findings led by (a) ROC Chart (and Index); (b) Score Ranking (or Accuracy Score); (c) Fit Statistics; (or Classification report) and (4) Output.

(a) ROC Chart (and Index)

In [66]:

ROC

```

y_pred_proba_dt = dt_model.predict_proba(X_test)
y_pred_proba_log_reg = log_reg_model.predict_proba(X_test)
y_pred_proba_nn = nn_model.predict_proba(X_test)
y_pred_proba_ensemble = voting.predict_proba(X_test_log)

roc_index_dt = roc_auc_score(y_test, y_pred_proba_dt[:, 1])
roc_index_log_reg = roc_auc_score(y_test, y_pred_proba_log_reg[:, 1])
roc_index_nn = roc_auc_score(y_test, y_pred_proba_nn[:, 1])
roc_index_ensemble = roc_auc_score(y_test_log, y_pred_proba_ensemble[:, 1])

print("ROC index on test for DT:", roc_index_dt)
print("ROC index on test for logistic regression:", roc_index_log_reg)
print("ROC index on test for NN:", roc_index_nn)
print("ROC index on voting classifier:", roc_index_ensemble)

fpr_dt, tpr_dt, thresholds_dt = roc_curve(y_test, y_pred_proba_dt[:,1])
fpr_log_reg, tpr_log_reg, thresholds_log_reg = roc_curve(y_test, y_pred_proba_log_reg[:,1])
fpr_nn, tpr_nn, thresholds_nn = roc_curve(y_test, y_pred_proba_nn[:,1])
fpr_ensemble, tpr_ensemble, thresholds_ensemble = roc_curve(y_test, y_pred_proba_ensemble[:,1])

plt.plot(fpr_dt, tpr_dt, label='ROC Curve for DT {:.3f}'.format(roc_index_dt), color='red', lw=0.5)
plt.plot(fpr_log_reg, tpr_log_reg, label='ROC Curve for Log reg {:.3f}'.format(roc_index_log_reg), color='green', lw=0.5)
plt.plot(fpr_nn, tpr_nn, label='ROC Curve for NN {:.3f}'.format(roc_index_nn), color='darkorange', lw=0.5)
plt.plot(fpr_ensemble, tpr_ensemble, label='ROC Curve for Ensemble {:.3f}'.format(roc_index_ensemble), color='darkorange', lw=0.5)

plt.plot([0, 1], [0, 1], color='navy', lw=0.5, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.0])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver operating characteristic example')
plt.legend(loc="lower right")
plt.show()


```

ROC index on test for DT: 0.6157864609264042

ROC index on test for logistic regression: 0.4997357932951725

ROC index on test for NN: 0.49995384047267355

ROC index on voting classifier: 0.7489632873881283



(b) Score Ranking (or Accuracy Score)

In [67]:

```
print("Accuracy score on test for DT:", accuracy_score(y_test_log, y_pred_dt))
print("Accuracy score on test for Logistic Regression:", accuracy_score(y_test_log, y_pred_log_reg))
print("Accuracy score on test for NN:", accuracy_score(y_test_log, y_pred_nn))
print("Accuracy score on test for Ensemble:", accuracy_score(y_test_log, y_pred_ensemble))
```

Accuracy score on test for DT: 0.8751105038977739

Accuracy score on test for Logistic Regression: 0.7545607972353934

Accuracy score on test for NN: 0.8350880012858636

Accuracy score on test for Ensemble: 0.8535722896407619

(c) Classification report

In [68]:

```
print("Report for DT: \n",classification_report(y_test_log, y_pred_dt))
print("\nReport for Logistic Regression: \n",classification_report(y_test_log, y
_pred_log_reg))
print("\nReport for NN: \n",classification_report(y_test_log, y_pred_nn))
print("\nReport for Ensemble: \n",classification_report(y_test_log, y_pred_ensem
ble))
```

Report for DT:

	precision	recall	f1-score	support
0	0.88	1.00	0.93	10832
1	0.80	0.05	0.09	1611
micro avg	0.88	0.88	0.88	12443
macro avg	0.84	0.52	0.51	12443
weighted avg	0.87	0.88	0.82	12443

Report for Logistic Regression:

	precision	recall	f1-score	support
0	0.93	0.78	0.85	10832
1	0.29	0.60	0.39	1611
micro avg	0.75	0.75	0.75	12443
macro avg	0.61	0.69	0.62	12443
weighted avg	0.85	0.75	0.79	12443

Report for NN:

	precision	recall	f1-score	support
0	0.90	0.91	0.91	10832
1	0.36	0.35	0.36	1611
micro avg	0.84	0.84	0.84	12443
macro avg	0.63	0.63	0.63	12443
weighted avg	0.83	0.84	0.83	12443

Report for Ensemble:

	precision	recall	f1-score	support
0	0.91	0.93	0.92	10832
1	0.43	0.37	0.40	1611
micro avg	0.85	0.85	0.85	12443
macro avg	0.67	0.65	0.66	12443
weighted avg	0.85	0.85	0.85	12443

(d) Output

In []:

b. Do all the models agree on the cars characteristics? How do they vary?

In []:

Task 6. Final Remarks: Decision Making

1. Finally, based on all models and analysis, is there

2. Can you summarise positives and negatives of each predictive modelling method based on this analysis?

3. How the outcome of this study can be used by decision makers?

In []:

In []:

In []: